



UNIVERSITY OF TRIESTE

Department of Mathematics, Informatics and Geosciences

Bachelor's Degree in Artificial Intelligence and Data Analytics

BACHELOR'S THESIS

**Cross-View Association of Identical Parts Across
Non-Overlapping 3D Scans: Identifiability,
Shift Ambiguity and Evaluation**

Supervisor:

Prof. Alejandro Rodriguez Garcia

Candidate:

Sabrina Ciccolo

Academic Year 2025/2026

Contents

Contents	iii
List of Figures	v
List of Tables	vii
1 Introduction	1
1.1 The problem	1
1.2 Related work	3
Bibliography	5

List of Figures

1.1	Dual-scanner system	2
1.2	The shift ambiguity (schematic)	2

List of Tables

Chapter 1

Introduction

1.1 The problem

Industrial handling is increasingly automated. Repetitive and heavy tasks such as lifting or sorting parts, once assigned to human operators, are now delegated to robots to satisfy the cycle time required by the production line.

For a robot to grasp an object, it is necessary to determine the object's pose, which dedicated sensors measure as a position and an orientation relative to the robot's own coordinate frame.

For a single compact piece, one 3D sensor is usually sufficient, as its viewpoint can cover the entire object. On the other hand, a very long part cannot be captured from a single viewpoint, making a single scanner insufficient. The layout therefore requires two sensors, one for each end. The part is gripped precisely at these ends, whose faces also determine the 6D pose of a straight extruded profile up to the symmetry of the cross-section.

The setup consists of two structured-light 3D scanners, positioned at either end of a conveyor carrying bundles of six-metre metal bars of identical profile.

Each device is mounted on a dedicated carriage at one end of the bundle, facing the end faces of the bars rather than their length. The two carriages are separated by several metres along the axis of the bar. The part of the bundle that each scanner covers is often narrower than the bundle itself, so a single view does not always contain every bar. Because the two sensors face opposite ends, their respective viewpoints neither overlap nor capture any common surface. Finally, a dual-arm manipulator picks up each bar by gripping it at both ends. Figure 1.1 illustrates this setup.

The task is therefore to determine which detection from one scanner and which from the other belong to the same physical object. In this specific case, data association cannot be resolved using conventional techniques due to the absence of typical geometric or visual features.

Absence of epipolar constraint. The lack of overlap between the two views means that the scanners never see the same surface.

Visual appearance invariance. Because the bars originate from the same extrusion batch, end faces belonging to the same profile family (i.e., bars sharing the same cross-section) are identical up to saw marks.




Figure 1.1: Dual-scanner system. Top view of the system. Each scanner is positioned at one end of the bundle and detects only the end faces of the bars facing it. In the figure, these faces are highlighted at both ends of each bar. The viewpoints of the two scanners do not share any surface and the bundle passes through both. Neither axis is to scale.

Relative position invariance. The bars are packed edge to edge at an almost uniform pitch, so the scanner detects a sequence of end faces that remains visually invariant even if the bundle shifts by one position.

Conveyor motion between acquisitions. Because the conveyor belt advances between the two acquisitions, the two scan windows can cut the bundle at different bars, in which case an association that pairs the detections in their order is shifted by one or more bars.

This shift ambiguity is illustrated in fig. 1.2. In each view, a different end bar remains excluded because each view often covers less than the whole bundle and the bundle moved between the two acquisitions. A correct association should record two unpaired detections, whereas a complete sequential association would lead to a systematic error.




Figure 1.2: The shift ambiguity (schematic). Each scanner covers a different part of the bundle, so a different bar is truncated in each acquisition. A correspondence shifted by one pitch would incorrectly pair all detections, whereas the correct correspondence leaves the two truncated detections unpaired. Letters identify the physical bar and hollow circles mark the unpaired detections.

This condition is not an edge case because in many scenes of the dataset some bars are

scanned at only one end. The objective is to determine what portion of the correspondence is actually recoverable and to define how unpaired detections are handled.

The scans used in this thesis were provided by beanTech, the company that developed the vision software of the rig for the manufacturer that operates it. At the time of this work the rig was still being commissioned and only a limited number of scenes was available. This setting favours methods that require no training data, such as the classical detector and the pairing methods studied here, since the detector needs only the drawing of the bar cross-section.

1.2 Related work

To the best of our knowledge, no method in the literature addresses this specific case. Cross-view association is a widely studied problem, but existing approaches rely on assumptions that do not hold in this setting, such as the presence of a distinctive visual appearance or the geometric overlap of fields of view. Four main lines of research can nevertheless be identified, each resting on assumptions whose absence in this setting has specific consequences.

Geometric correspondence and assignment. The first line of research assumes the presence of surfaces observed in common. Methods based on stereo matching and multi-view correspondence, starting from dynamic programming along scanlines Ohta et al. 1985 [7], function precisely because the same surface appears in multiple images. The same principle applies to systems that associate identical instances detected by different cameras using geometry alone. These systems calculate the epipolar distance for each pair of detections, model it as a graph cost and recover the groups by clustering Sun et al. 2025 [11]. In the absence of visual overlap, however, the epipolar constraint no longer applies. The task thus reduces to a pure assignment problem, in which the cost matrix must be derived directly from the scene layout. The Hungarian algorithm solves minimum-cost bipartite matching Kuhn 1955 [3] without imposing order constraints on correspondences. An assignment can also be constrained to preserve the order of the elements, which is an explicit modelling assumption and does not follow from epipolar geometry. The monotonic alignment compared in this thesis (??) imposes such a constraint on the order of the bars.

Discriminative appearance. The second line of research leverages visual appearance to discriminate one object from another. Benchmarks for cross-camera instance association, such as MessyTable Cai et al. 2020 [1], are designed for objects showing appearance differences between instances, even minimal ones. According to their analysis, reliable association requires appearance, context and geometry together. The application most closely related to the case under consideration is found in the timber industry, where images of the same log end face taken at different times are matched as a biometric fingerprint, either with texture features that describe the annual ring pattern Schraml et al. 2015 [9] or with features learned by a deep network Xie et al. 2025 [13]. This approach is effective because each end face presents a unique surface, unlike a metallic extrusion whose cross-section is geometrically and visually identical to adjacent ones.

Identifiability and observability. The third line of research concerns the information content of the data themselves. In the absence of overlap and characterising visual attributes, the task reduces to an unlabelled sensing problem Unnikrishnan et al. 2018 [12]: observations are known, but the correspondence matrix is unknown. The fundamental question lies in determining what can actually be inferred from them. The literature addresses the topic along two directions: defining which configurations are reconstructible and formulating a converse that bounds what no algorithm can recover. The same critical issue emerges

in self-calibration for system parameter estimation, where it is addressed by studying the observability of the information matrix Maye et al. 2016 [5] and identifying unobservable directions in the parameter space Nobre et al. 2019 [6].

Association performance evaluation. The fourth line of research concerns association evaluation, which requires metrics distinct from those used for pure detection quality. Several studies on multi-target multi-camera systems highlight that single-frame metrics do not provide information on identity consistency across space and time. For this reason, identity-focused precision and recall metrics were introduced Ristani et al. 2016 [8], followed later by the HOTA metric, which prevents an aggregate index from masking the source of error between detection and data association Luiten et al. 2021 [4]. The BOP benchmark applies an analogous principle to 6D pose estimation, defining correctness in relation to the symmetry group of the object rather than relative to a canonical pose Hodaň et al. 2020 [2]. Finally, concurrent work shows that pairwise ranking metrics and one-to-one assignment metrics can disagree, which means that a model may achieve a higher score in ranking metrics without a corresponding improvement in its actual assignment capability Shelukhan et al. 2026 [10].

The gap addressed by this work lies between identifiability and evaluation. For an acquisition setup without visual overlap that observes identical elements at a regular spacing, it has not yet been formalised which part of the matching problem can be resolved in principle and which part is intrinsically undetermined. Consequently, reading an overall accuracy metric requires separating correctable algorithmic errors from structural limitations of the physical setup.

Bibliography

- [1] Z. Cai et al. “MessyTable: Instance Association in Multiple Camera Views”. In: *Computer Vision – ECCV 2020*. Lecture Notes in Computer Science. 2020, pp. 1–16. DOI: 10.1007/978-3-030-58621-8_1.
- [2] T. Hodaň et al. “BOP Challenge 2020 on 6D Object Localization”. In: *Computer Vision – ECCV 2020 Workshops*. 2020. arXiv: 2009.07378 [cs.CV].
- [3] H. W. Kuhn. “The Hungarian method for the assignment problem”. In: *Naval Research Logistics Quarterly* 2.1-2 (1955), pp. 83–97. DOI: 10.1002/nav.3800020109.
- [4] J. Luiten et al. “HOTA: A Higher Order Metric for Evaluating Multi-Object Tracking”. In: *International Journal of Computer Vision* 129.2 (2021), pp. 548–578. DOI: 10.1007/s11263-020-01375-2.
- [5] J. Maye et al. “Online self-calibration for robotic systems”. In: *The International Journal of Robotics Research* 35.4 (2016), pp. 357–380. DOI: 10.1177/0278364915596232.
- [6] F. Nobre and C. R. Heckman. *FastCal: Robust Online Self-Calibration for Robotic Systems*. 2019. arXiv: 1902.10585 [cs.R0].
- [7] Y. Ohta and T. Kanade. “Stereo by Intra- and Inter-Scanline Search Using Dynamic Programming”. In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* PAMI-7.2 (1985), pp. 139–154. DOI: 10.1109/TPAMI.1985.4767639.
- [8] E. Ristani et al. “Performance Measures and a Data Set for Multi-Target, Multi-Camera Tracking”. In: *Computer Vision – ECCV 2016 Workshops*. 2016. arXiv: 1609.01775 [cs.CV].
- [9] R. Schraml et al. “Towards the applicability of biometric wood log traceability using digital log end images”. In: *Computers and Electronics in Agriculture* 119 (2015), pp. 112–122. DOI: 10.1016/j.compag.2015.10.003.
- [10] M. Shelukhan et al. *Ranking vs. Assignment: The Metric Mismatch in Multi-View Object Association*. 2026. arXiv: 2606.02022 [cs.CV].
- [11] Y.-H. Sun et al. *C-DOG: Multi-View Multi-instance Feature Association Using Connected δ -Overlap Graphs*. 2025. arXiv: 2507.14095 [cs.CV].
- [12] J. Unnikrishnan, S. Haghighatshoar, and M. Vetterli. “Unlabeled Sensing With Random Linear Measurements”. In: *IEEE Transactions on Information Theory* 64.5 (2018), pp. 3237–3253. DOI: 10.1109/TIT.2018.2809002.

- [13] Y. Xie et al. “Log End Face Feature Extraction and Matching Method Based on Swin Transformer V2”. In: *Forests* 16.1 (2025), p. 124. DOI: 10.3390/f16010124.