



UNIVERSITY OF TRIESTE

Department of Mathematics, Informatics and Geosciences

Bachelor's Degree in Artificial Intelligence and Data Analytics

BACHELOR'S THESIS

**Cross-View Association of Identical Parts Across
Non-Overlapping 3D Scans: Identifiability,
Shift Ambiguity and Evaluation**

Supervisor:

Prof. Alejandro Rodriguez Garcia

Candidate:

Sabrina Ciccolo

Academic Year 2025/2026

Abstract

A robot that picks long metal bars from a conveyor belt must first locate them. In the industrial rig studied here, the bars form a bundle with a 3D scanner at each end, so locating a bar requires pairing the end faces seen by one scanner with those seen by the other but the two views share no surface and the bars are identical and almost evenly spaced. In addition, the scan window of each scanner is often narrower than the bundle, which moves between the two acquisitions, so one view can contain bars that the other does not see.

Existing methods for cross-view association rely on surfaces visible in both views or on objects that differ in appearance and do not appear suited to this case. Besides the accuracy of pairing methods, this thesis therefore studies which part of the correspondence the data can determine and how much ambiguity remains, combining an identifiability analysis, a synthetic twin of the rig and 108 real scenes.

Several pairing methods minimise a cost made of a geometric term, which measures how poorly the paired end faces align, and a counting term, which penalises every detection left unpaired. When the bars are almost evenly spaced, an association shifted by one bar has practically the same geometric term as the true correspondence and only the counting term tells the two apart, although this term is no longer enough if the two views contain partly different bars because the true correspondence then also leaves some detections unpaired.

This mechanism is the central result of the thesis. On real data, the most accurate method recovers 92.9% of the annotated correspondences but in every scene on which it fails completely, which always contains at least one detection without an annotated partner, it returns a shifted association that costs less than the true correspondence. The error therefore lies in the cost itself and not in how its minimum is found. Eliminating it would require information that distinguishes individual bars.

From the scans it is also possible to recover the separation between the two scanners along the bar axis but not their absolute position, which does not change the reconstructed bar length. Estimating the separation therefore removes the systematic error of about one metre that the offline calibration available to this work leaves in the reconstructed length. Finally, a shifted association forms almost as many pairs as the true correspondence but is wrong on every bar, so the thesis proposes an evaluation that checks whether the two end faces of each bar have been paired with each other.

Contents

Abstract	i
Contents	iii
List of Figures	v
List of Tables	vii
1 Introduction	1
1.1 The problem	1
1.2 Related work	3
1.3 Contributions	4
2 The data, the rig and the task	7
2.1 The dataset, annotations and evaluation metrics	7
2.2 Acquisition geometry	9
2.3 Detection	9
2.4 The association task and the methods compared	13
2.5 The shift family and identifiability	16
2.6 The synthetic twin and its limits	16
3 Identifiability of the carriage geometry	19
3.1 The length prior and its residual	19
3.2 Identifiability of the carriage positions	20
3.3 Recovering the carriage separation	22
3.4 Calibration in Z , transport in Y	24
3.5 Metric recovery and the combinatorial ambiguity	24
4 Shift ambiguity after self-calibration	25
4.1 Sensitivity of the length check to a one-bar shift	25
4.2 The ambiguity as a function of cross-view transport	26
4.3 Truncation and transport	28
4.4 Cost of the true correspondence on real scenes	29

5	Evaluation of the pairing methods	31
5.1	Counting and correspondence metrics on the same detections	31
5.2	Structure of the pairing errors	32
5.3	Correspondence accuracy of the pairing methods	34
5.4	Stability under a change of detector	35
5.5	Comparison with the original system’s pairing criterion	38
6	Conclusion	39
6.1	Summary of the results	39
6.2	Implications for pairing and evaluation	40
6.3	Limitations	41
6.4	Future work	41
	Bibliography	43

List of Figures

1.1	Dual-scanner system	2
1.2	The shift ambiguity (schematic)	2
2.1	Two views of one scene in the world frame	11
2.2	Detection	12
2.3	The same bundle seen from the two ends	13
2.4	The four cross-sections	14
2.5	Twin against real	17
3.1	Unobservable direction of the length prior	21
3.2	Cross-scene reproducibility of the carriage separation	23
4.1	The ambiguity sweep	27
5.1	Effect of a change of detector	36

List of Tables

2.1	Profile families of the dataset	7
2.2	Quantities of the rig and of the acquisitions	10
3.1	Bar length residual under the three calibrations	22
4.1	Length change caused by a one-bar shift	26
4.2	Geometric and counting terms of the objective	28
4.3	Truncation against transport	29
5.1	Counting and correspondence metrics of the pairing methods	32
5.2	Classification of the scenes by pairing outcome	33
5.3	Correspondence recall compared with the hybrid method	34
5.4	Stability of correspondence recall under a change of detector	37

Chapter 1

Introduction

1.1 The problem

Industrial handling is increasingly automated. Repetitive and heavy tasks such as lifting or sorting parts, once assigned to human operators, are now delegated to robots to satisfy the cycle time required by the production line.

For a robot to grasp an object, it is necessary to determine the object's pose, which dedicated sensors measure as a position and an orientation relative to the robot's own coordinate frame.

For a single compact piece, one 3D sensor is usually sufficient, as its viewpoint can cover the entire object. On the other hand, a very long part cannot be captured from a single viewpoint, making a single scanner insufficient. The layout therefore requires two sensors, one for each end. The part is gripped precisely at these ends, whose faces also determine the 6D pose of a straight extruded profile up to the symmetry of the cross-section.

The setup consists of two structured-light 3D scanners, positioned at either end of a conveyor carrying bundles of six-metre metal bars of identical profile.

Each device is mounted on a dedicated carriage at one end of the bundle, facing the end faces of the bars rather than their length. The two carriages are separated by several metres along the axis of the bar. The part of the bundle that each scanner covers is often narrower than the bundle itself, so a single view does not always contain every bar. Because the two sensors face opposite ends, their respective viewpoints neither overlap nor capture any common surface. Finally, a dual-arm manipulator picks up each bar by gripping it at both ends. Figure 1.1 illustrates this setup.

The task is therefore to determine which detection from one scanner and which from the other belong to the same physical object. In this specific case, data association cannot be resolved using conventional techniques due to the absence of typical geometric or visual features.

Absence of epipolar constraint. The lack of overlap between the two views means that the scanners never see the same surface.

Visual appearance invariance. Because the bars originate from the same extrusion batch, end faces belonging to the same profile family (i.e., bars sharing the same cross-section) are identical up to saw marks.




Figure 1.1: Dual-scanner system. Top view of the system. Each scanner is positioned at one end of the bundle and detects only the end faces of the bars facing it. In the figure, these faces are highlighted at both ends of each bar. The viewpoints of the two scanners do not share any surface and the bundle passes through both. Neither axis is to scale.

Relative position invariance. The bars are packed edge to edge at an almost uniform pitch, so the scanner detects a sequence of end faces that remains visually invariant even if the bundle shifts by one position.

Conveyor motion between acquisitions. Because the conveyor belt advances between the two acquisitions, the two scan windows can cut the bundle at different bars, in which case an association that pairs the detections in their order is shifted by one or more bars.

This shift ambiguity is illustrated in fig. 1.2. In each view, a different end bar remains excluded because each view often covers less than the whole bundle and the bundle moved between the two acquisitions. A correct association should record two unpaired detections, whereas a complete sequential association would lead to a systematic error.




Figure 1.2: The shift ambiguity (schematic). Each scanner covers a different part of the bundle, so a different bar is truncated in each acquisition. A correspondence shifted by one pitch would incorrectly pair all detections, whereas the correct correspondence leaves the two truncated detections unpaired. Letters identify the physical bar and hollow circles mark the unpaired detections.

This condition is not an edge case because in many scenes of the dataset some bars are scanned at only one end. The objective is to determine what portion of the correspondence is actually recoverable and to define how unpaired detections are handled.

The scans used in this thesis were provided by beanTech, the company that developed the vision software of the rig for the manufacturer that operates it. At the time of this work the rig was still being commissioned and only a limited number of scenes was available. This setting favours methods that require no training data, such as the classical detector and the pairing methods studied here, since the detector needs only the drawing of the bar cross-section.

1.2 Related work

To the best of our knowledge, no method in the literature addresses this specific case. Cross-view association is a widely studied problem, but existing approaches rely on assumptions that do not hold in this setting, such as the presence of a distinctive visual appearance or the geometric overlap of fields of view. Four main lines of research can nevertheless be identified, each resting on assumptions whose absence in this setting has specific consequences.

Geometric correspondence and assignment. The first line of research assumes the presence of surfaces observed in common. Methods based on stereo matching and multi-view correspondence, starting from dynamic programming along scanlines Ohta et al. 1985 [10], function precisely because the same surface appears in multiple images. The same principle applies to systems that associate identical instances detected by different cameras using geometry alone. These systems calculate the epipolar distance for each pair of detections, model it as a graph cost and recover the groups by clustering Sun et al. 2025 [14]. In the absence of visual overlap, however, the epipolar constraint no longer applies. The task thus reduces to a pure assignment problem, in which the cost matrix must be derived directly from the scene layout. The Hungarian algorithm solves minimum-cost bipartite matching Kuhn 1955 [5] without imposing order constraints on correspondences. An assignment can also be constrained to preserve the order of the elements, which is an explicit modelling assumption and does not follow from epipolar geometry. The monotonic alignment compared in this thesis (section 2.4) imposes such a constraint on the order of the bars.

Discriminative appearance. The second line of research leverages visual appearance to discriminate one object from another. Benchmarks for cross-camera instance association, such as MessyTable Cai et al. 2020 [1], are designed for objects showing appearance differences between instances, even minimal ones. According to their analysis, reliable association requires appearance, context and geometry together. The application most closely related to the case under consideration is found in the timber industry, where images of the same log end face taken at different times are matched as a biometric fingerprint, either with texture features that describe the annual ring pattern Schraml et al. 2015 [12] or with features learned by a deep network Xie et al. 2025 [16]. This approach is effective because each end face presents a unique surface, unlike a metallic extrusion whose cross-section is geometrically and visually identical to adjacent ones.

Identifiability and observability. The third line of research concerns the information content of the data themselves. In the absence of overlap and characterising visual attributes, the task reduces to an unlabelled sensing problem Unnikrishnan et al. 2018 [15]: observations are known, but the correspondence matrix is unknown. The fundamental question lies in determining what can actually be inferred from them. The literature addresses the topic along two directions: defining which configurations are reconstructible and formulat-

ing a converse that bounds what no algorithm can recover. The same critical issue emerges in self-calibration for system parameter estimation, where it is addressed by studying the observability of the information matrix Maye et al. 2016 [8] and identifying unobservable directions in the parameter space Nobre et al. 2019 [9].

Association performance evaluation. The fourth line of research concerns association evaluation, which requires metrics distinct from those used for pure detection quality. Several studies on multi-target multi-camera systems highlight that single-frame metrics do not provide information on identity consistency across space and time. For this reason, identity-focused precision and recall metrics were introduced Ristani et al. 2016 [11], followed later by the HOTA metric, which prevents an aggregate index from masking the source of error between detection and data association Luiten et al. 2021 [7]. The BOP benchmark applies an analogous principle to 6D pose estimation, defining correctness in relation to the symmetry group of the object rather than relative to a canonical pose Hodaň et al. 2020 [3]. Finally, concurrent work shows that pairwise ranking metrics and one-to-one assignment metrics can disagree, which means that a model may achieve a higher score in ranking metrics without a corresponding improvement in its actual assignment capability Shelukhan et al. 2026 [13].

The gap addressed by this work lies between identifiability and evaluation. For an acquisition setup without visual overlap that observes identical elements at a regular spacing, it has not yet been formalised which part of the matching problem can be resolved in principle and which part is intrinsically undetermined. Consequently, reading an overall accuracy metric requires separating correctable algorithmic errors from structural limitations of the physical setup.

1.3 Contributions

This thesis determines what the data of two non-overlapping views can recover. The geometry of the rig is recovered except for the absolute position of the scanners, whereas when the bars are evenly spaced the geometry of the scene determines the correspondence between the detections only up to a shift by a whole number of bars. The pairing methods are then evaluated on the basis of this distinction rather than through a single aggregate accuracy figure. The contributions of this work are the following.

1. **The mechanism of the shift ambiguity.** Besides the true association, the data are compatible with a family of associations shifted by a whole number of bars along the bundle. When the bars are evenly spaced, a shifted association has practically the same geometric cost as the true one, because the translation estimated for the scene changes by the same shift and the distances between paired faces stay the same. The two associations differ only in the number of detections left unpaired, which is lower for the true association when both views contain the same bars. For this reason the ambiguity depends on whether the scan window truncates the two views differently and not on how far the bundle moves between the two acquisitions (sections 2.5, 4.2 and 4.3).
2. **Confirmation on real data.** The errors of the most accurate pairing method on the real scenes follow this mechanism. Every scene on which the method fails completely is shifted as a whole and contains at least one detection without an annotated partner. In each of these scenes the wrong association also has a lower cost than the annotated one under the objective that the method minimises, so these errors depend on the

information contained in the detections rather than on the search for the minimum (section 5.2).

3. **Identifiability of the rig geometry.** Prior knowledge of the nominal bar length determines the separation between the two carriages but not their absolute position, which the scene content does not constrain. Alternating calibration and assignment recovers the separation and removes an error of about one metre along the bar axis (sections 3.2 and 3.3). The length check of the original system cannot detect a shifted association, because pairing the ends of two adjacent bars changes the reconstructed length by far less than the tolerance of the check (section 4.1).
4. **The hybrid method.** The pairing method proposed in this thesis applies the order constraint of the monotonic alignment within each stacking level and admits a bar crossing its neighbours only when the detections support it (section 2.4). With both of the detectors compared in this thesis it is the most accurate of the methods tested, although on this dataset its advantage over pairing the detections by their order alone is not statistically significant (sections 5.3 and 5.4).
5. **Evaluation of the pairing methods.** The counting metric checks only how many pairs are assigned and therefore does not distinguish a correct association from a shifted or swapped one. This thesis relies instead on a correspondence annotation linked to the detections by position, which makes it possible to score the output of any detector against the same ground truth (sections 2.1 and 5.1).
6. **Development and audit of a synthetic twin.** The synthetic twin, developed to simulate scene acquisitions, replicates detection counts and detector accuracy but does not model how far the bar ends stagger along the bar axis. This restricts the scope of any validation conducted on it to the transverse plane alone (section 2.6).

Chapter 2

The data, the rig and the task

Apart from the controlled experiments on the synthetic twin (section 2.6), all results in this thesis come from real acquisitions. Every comparison is between pairing methods that operate on the same detections and return cross-view pairs. A detection is an end face identified in a single view, whereas a pairing method is the procedure that takes the detections from both views as input and returns a set of pairs connecting them.

2.1 The dataset, annotations and evaluation metrics

A scene consists of a bundle of bars acquired once from each end, yielding two point clouds. The dataset comprises 108 scenes and 216 clouds, classified by profile family (table 2.1), with each scene associated with a manually assigned scenario category.

Table 2.1: Profile families of the dataset. For each family, the shape of its cross-section as given by the CAD drawing, the number of scenes and the number of annotated faces. The L-profile appears twice because it was acquired in two separate campaigns, so the five families share four distinct cross-sections. The five families total 108 scenes and 935 faces.

profile family	cross-section	scenes	annotated faces
complex profile	complex	4	15
heavy profile	rectangular, 80×95 mm	11	68
L-profile	asymmetric L-shape	40	274
L-profile, second acquisition	asymmetric L-shape	24	223
square profile	square, 40×40 mm	29	355

The second column indicates the cross-section shape, obtained from the CAD drawing of the profile. Although the dataset consists of five families, there are only four distinct cross-sections because the L-profile appears twice due to two separate acquisition campaigns.

The categories describe the arrangement of the bars in the bundle: single (56 scenes), overlapping (22), standing (15), crossed (9) and tilted (6). Across all scenes, a single side features a median of 4 faces and a maximum of 29.

Truncation, where a bar is visible to only one scanner, affects 30 of the 108 scenes and leaves 49 of the 935 faces without a counterpart in the other view. Consequently, the assumption that every bar is visible to both scanners fails in roughly one out of four scenes.

Two types of annotations have been defined.

The first is a count annotation, indicating how many end faces are visible for each side. It was made first, by hand, on images generated from the point clouds and independently of any detector.

The second is a correspondence annotation that, for each scene, provides a list of claims: inter-view pairs, each equipped with an anchor point on both faces, taken from the centroid of the face detected by the candidate-pool detector and confirmed by the annotator. In total, this annotation contains 407 claims, with at least one in every scene. The anchors allow each claim to be associated with the detections of the detector being evaluated, verifying whether the detector successfully produced that pair. A claim is resolved for a detector when both of its anchors can be linked to a detection of that detector. Where a scene needs it, the annotation also carries a note recording, for example, which profile is truncated in one view or not confirmed by the detector.

The two annotations support different metrics. A counting metric checks only how many pairs a method produces, so the criterion is satisfied whenever the number of pairs is correct, even if all pairs are incorrect. For example, if two true pairs are swapped, the method still produces two pairs and the counting metric counts them as two true positives. The same happens when one view truncates a bar that the other contains, in which case an association shifted by one position along the row keeps the annotated pair count. Only the correspondence annotation indicates which pairs are correct and can therefore detect a swap. Work on multi-camera tracking faced a similar limitation and resolved it with identity-based precision and recall metrics Ristani et al. 2016 [11].

Four metrics are evaluated:

- **Counting precision, recall and F1:** compare the number of faces or pairs produced by a method against the count annotation, both scene-by-scene and aggregated. A true positive is defined as the smaller of the two numbers, calculated per side for faces and per scene for pairs.
- **Correspondence recall:** the proportion of claims correctly reproduced by a method (i.e., with correct bar identity), evaluated over the set of resolved claims. Restricting the score to resolved claims prevents a missing detection from being counted as a pairing error, following the distinction formalised by HOTA for tracking Luiten et al. 2021 [7].
- **IPAA-1.0** (Image-Pair Association Accuracy at 1.0): the proportion of scenes correctly resolved in their entirety (referred to as IPAA-100 in MessyTable Cai et al. 2020 [1]). While aggregated recall gives greater weight to scenes with many associations, IPAA-1.0 weighs every scene equally. Furthermore, it aligns most closely with a line operator’s experience, where an entire bundle is processed as a single unit.
- **Inversions:** the count of pairs within the same row that lie side-by-side at the same height in both views but exhibit an inverted Y-order between them. These capture the swaps that counting metrics cannot detect by design.

In this thesis, all confidence intervals are calculated by resampling bundles rather than pairs. Pairs within the same scene share the same bundle, transport estimate and assignment, so they are not independent samples and an error tends to affect the whole scene rather than individual pairs. Some scenes are also repeat captures of the same stationary bundle, whose point clouds coincide within 2 mm on almost every point. These scenes are grouped and resampled together, so the 108 scenes of the dataset form 96 independent units, called bundles

in the rest of the thesis. For correspondence metrics all of them are resampled, since each scene contains at least one claim.

All results reported here are based on the candidate-pool detector (section 2.3). The anchors were defined from its output, so every claim corresponds to one of its detections and all 407 claims can be scored.

The alternative detector, based on two-pass suppression (section 2.3), resolves 382 of the 407 claims and re-associates faces with a median error of 0.96 mm, compared to 0.00 mm for the candidate-pool detector. The comparison between the two detectors therefore measures how robust each method is to a change of detector rather than which detector is more accurate.

2.2 Acquisition geometry

Two scanners are placed at opposite ends of a conveyor belt, each mounted on its own linear carriage. The global reference frame (in millimetres) is defined relative to the robot: the X-axis aligns with the bars and carriage motion, the Y-axis indicates the conveyor feed direction and the Z-axis represents height. Calibration relates each scanner’s local frame to the global frame via a rigid transformation $p_{\text{world}} = R p_{\text{scanner}} + t$. After calibration, the point clouds are orthographically projected onto the YZ plane (feed and height), which lies nearly frontal to the end faces.

Two calibrations are available. The offline calibration, defined by a single rotation and translation knot per scanner, yields a constant rigid transformation valid exclusively at the carriage position where knots were acquired. By contrast, the original system calibration uses ten knots per scanner across roughly 2.1 m of travel, interpolating between them with a centripetal Catmull-Rom spline for translation and spherical cubic interpolation (SQUAD) for rotation, with values clamped at the boundaries. Applying the original calibration would require the carriage position of every scan, which is absent from the dataset, so the constant offline calibration was used instead. Consequently, accuracy is limited to the single acquired position, introducing an error along the X coordinate, the axis of carriage motion, that is subsequently measured and compensated (section 3.3). Figure 2.1 shows the two clouds of one scene in the world frame.

Several quantities in table 2.2, along with certain detection constants, originate from the source code of the original system installed on the machine.

2.3 Detection

The points that survive the radial filter are projected onto the YZ plane and a look-up table records which 3D points correspond to each pixel, enabling a 2D match to retrieve its 3D support. The profile cross-section is rasterised from the CAD drawing (DXF file) at 2.5 px/mm and slid over the image at 2° intervals across the full rotation. Every local cross-correlation maximum reaching the lost-face threshold of 0.2 constitutes a candidate (the threshold is defined relative to the template’s self-correlation peak).

The two detectors used in this work share this initial stage but differ in the criterion used to confirm candidates.

The two-pass suppression detector decides on the 2D image: it confirms a candidate when the peak reaches the match threshold of 0.4 and the mask overlap exceeds 0.65, suppresses

Table 2.2: Quantities of the rig and of the acquisitions. For each quantity, its value and how it was obtained. Lengths are in millimetres in the world frame, where X runs along the bars, Y along the conveyor feed and Z in height.

Quantity	Value	How it was obtained
Standoff	656 mm, interquartile range (IQR) [641, 669] mm	Median radial distance between the scanner origin and the retained points across 216 clouds.
Angle of incidence on the end face	22.0°	Extracted from the production calibration rotation matrices. Stable within $\pm 0.35^\circ$ over roughly 2.1 m of carriage travel.
Angular envelope	$[-19.3^\circ, +17.4^\circ]$, union of both sides	Field of view width of each scanner along its tangential direction, within its local frame. The lower limit repeats within 0.06° across 104 of 108 scenes on both sides, indicating a hardware limitation.
Scan window	median 258 mm, maximum 414 mm	Calculated between the 0.5th and 99.5th percentiles of points retained after the 850 mm radial filter. Prior to filtering, the median extent is 479 mm. The effective window depends jointly on this radial filter and the device’s angular span.
Bar pitch	[38.6, 118.4] mm	Range of median distances between adjacent bars, calculated for each of the five families on the correspondence anchors.
Cross-view transport	median $ \Delta Y $ 32.1 mm	Calculated per scene as the absolute value of the median ΔY between the anchor points of the two views. The reported value represents the median of these values across all 108 scenes.
Nominal bar length	6005 mm	Value set by the operator in the original system, independent of the profile family.
Distance between carriages	5403.3 mm	Estimated via self-calibration (section 3.3) rather than direct measurement. It differs by 1.3 mm from the documented separation of 2701 mm per side (5402 mm total).
Effective resolution	2.5 px/mm	Product of the rendering scale (10 px/mm) and the downscaling factor (0.25).




Figure 2.1: Two views of one scene in the world frame. The two point clouds of a single scene in the global reference frame. The radial filter, applied before calibration, discards any point beyond 850 mm from its scanner, which isolates the end faces from the machine and the conveyor belt and leaves 15,037 of the 298,061 points on the first side. Nothing is ever acquired between the two clouds, so the bar is never fully scanned. The scene belongs to the L-profile family.

duplicates at an intersection over union (IoU) of 0.2, clears the confirmed footprints from the image and repeats the sweep so that a hidden face behind a stronger neighbour can emerge in the second pass.

The candidate-pool detector instead decides based on the 3D points for which candidates compete. It retains every candidate above 0.2 without angle limits or bounding box overlap suppression, assigning to each candidate the 3D points within its footprint at its depth. Finally, a single global selection picks the subset explaining the most points against a fixed per-instance cost, with the constraint that no two selected candidates claim the same point. The candidate-pool detector is the one used in the remainder of this work, while two-pass suppression was retained to repeat the comparison between pairing methods on a different underlying detector. Figure 2.2 shows how the two detectors interpret the same image.

For each detected face, both methods return the 3D footprint points by masking the crop with the template shape and tracing back through the look-up table. Extracting the centroid and normal of the fitted plane via total least squares yields inputs for the pairing gates, combined with the matching angle retained from the 2° quantised sweep.

In this detection phase, the resolution, the 2° step size and the correlation thresholds come from the original system installed with the machine. After identifying a face and retrieving its 3D points, that system executed an extra step: it re-projected the face onto its own plane, estimated using RANSAC Fischler et al. 1981 [2], and repeated the angle search on that view with a finer step. It thus obtained a refined angle and a planar reference for each face. The detection pipeline used in this thesis does not perform this step. A more precise angle would tighten the consistency checks between the ends of the same bar without helping to decide which end belongs to which bar, since identical parallel bars share the same angle.

Across all 108 scenes, the candidate-pool detector achieves a precision of 0.986, a recall of 0.955 and an F1 score of 0.970 against per-side counts. A subsequent recovery stage increases recall to 0.967 and F1 to 0.976. Operating on unpaired faces, this stage estimates




Figure 2.2: Detection. On the left, the projection on which both detectors operate after the radial filter and the template sweep. In the centre and on the right, the interpretation of the same image by each detector: confirmed detections are enclosed in coloured bounding boxes, while unconfirmed candidates appear in dashed grey boxes. In this case, both confirm the same five faces, but differ in the amount of residual unconfirmed evidence: four candidates versus seven.

their position in the opposing view using the cross-view translation, searches for a local match and accepts it only if the 3D footprint exhibits valid end-face characteristics. Approximately 20% of attempts are accepted.

As a baseline for the classical detector, a small convolutional network was trained on the same annotations and with the same code. It predicts face centres as peaks of a Gaussian heatmap, with a second head recovering their offset within the output cell. This formulation follows the centre-point detection proposed in Zhou et al. 2019 [17], whose Gaussian target and loss function derive from Law et al. 2018 [6]. Detection labels give only face centres, so a detector predicting bounding box extent and orientation would lack data to learn those parameters. Direct centre prediction instead keeps both model branches focused on per-side counts and the position of each face.

The results of the learned model depend on how the data is split. In the scene-disjoint split, where the training and test scenes differ but contain the same cross-sections, the model achieves a test F1 of 0.902. In the family-disjoint split, the test family is never seen during training and the two L-profile acquisitions are held out together, yielding an F1 of 0.425. The drop from 0.902 to 0.425 shows that the model recognises previously seen sections well, but fails to generalise to novel geometry. Both values are averages across two training runs, in each of which the model sees only a subset of the annotations, and should be compared with the classical detector’s 0.970. A paired bootstrap over bundles on one of the two runs gives, in both splits, a confidence interval for the difference that lies entirely below zero, so the learned model is significantly less accurate than the classical detector.

The discrepancy between the two splits shows that the model learns the four cross-sections present in the dataset, but does not develop a general end-face detector. Given a DXF drawing, the classical detector requires no training data and therefore always operates in the family-disjoint regime. The price is inference time, which is 2.53 s per scene side against 18 to 24 ms for the learned model. The two F1 values are lower bounds, because training was stopped before the test F1 had converged under a CPU-only schedule. The qualitative conclusion does not depend on them, since the model does not generalise to an unseen cross-section.

2.4 The association task and the methods compared

The association task starts from the faces detected in the two views and determines which faces, one per view, represent the two ends of the same bar. The inputs are therefore the detections from the two views and the output is a set of pairs between detections. Figure 2.3 shows an example.




Figure 2.3: The same bundle seen from the two ends. YZ projection on which the two detectors operate for the same scene. The two views are drawn on the same Y-axis of the world reference frame, making the face positions directly comparable. The faces are identical up to the saw marks but the two views do not coincide along Y because the conveyor belt moves between the two acquisitions. The scene belongs to the L-profile family.

The faces detected in a view are grouped into stacking levels, i.e., the horizontal rows of the bundle. A face is one level above another when the Y coordinate of its centroid falls within the Y interval of the first and its centroid's Z coordinate exceeds the maximum Z of the second. This relationship is iterated until stabilisation and the faces in each level are sorted by decreasing Y. The grouping is relative to the bundle because no belt plane measurement is available and the original system provides no nominal value for it.

The task is a partial association between the detections in the two views. Each detection includes a YZ centroid, a coarse matching angle, the estimated face normal and the stacking level to which it belongs. Given the detections $D_1 = \{d_1, \dots, d_m\}$ and $D_2 = \{e_1, \dots, e_n\}$ of the two sides, an association is a partial matching $\pi \subseteq D_1 \times D_2$ in which each detection appears at most once. Unpaired detections are allowed and, in this dataset, necessary because 49 faces have no counterpart.

All association methods compared in this thesis share a single cost model:

$$c(i, j) = \frac{|\Delta\theta(i, j)|}{\tau_\theta} + \min\left(\frac{\|\Delta_{YZ}(i, j)\|}{\tau_d}, \kappa\right) + \lambda_L |\text{level}(i) - \text{level}(j)| + \lambda_Z [\text{outlier}(i) \neq \text{outlier}(j)] \quad (2.1)$$

The cost consists of four terms. The first term, $\Delta\theta$, measures the orientation difference between the two faces. The cross-section can be symmetric, so two orientations differing by a full period are considered equal: 90° for a square section, 180° for a rectangular section and 360° for an asymmetric section. The period is measured from the section contour and is not declared per family, in line with symmetry handling in BOP 6D benchmarks Hodañ et al. 2020 [3]. Figure 2.4 shows the four cross-sections with their corresponding measured periods.




Figure 2.4: The four cross-sections. The five families of the dataset include four distinct cross-sections, as the L-profile was acquired twice. Each section is read from the technical drawing DXF through the same loader used by the detector’s template renderer. For each section, the bounding box, the bar pitch measured for each family using it and the rotational symmetry period adopted by the cost model are shown. The L-profile panel shows two pitch values, one per acquisition. The square section has a 90° period, the rectangular section 180° , while the L-shaped and complex sections have a 360° period.

The second term, Δ_{YZ} , is the distance between the centroids in the projection plane, saturated at κ . Saturation is necessary because the translation between the views is estimated and never exact, so the distance must remain a soft prior that still allows a transverse bar to be paired.

The third term penalises pairs belonging to different stacking levels. The fourth term penalises pairs in which one face is a height outlier and the other is not, where a face counts as an outlier when its centroid Z sits more than 8 mm from the median height of its own view. This term is kept separate from the distance term because it still discriminates after the distance has saturated. The scaling factors make a cost of 1 equivalent to the tolerance threshold: $\tau_\theta = 2.5^\circ$, $\tau_d = 45$ mm, $\lambda_L = \lambda_Z = 2$.

Two constants vary depending on the method. A method that does not estimate the cross-view translation uses $\kappa = 3$ and keeps a pair only when its cost is below 6, the cost of leaving both of its faces unpaired. Conversely, methods that estimate the translation increase these values to 4.5 and 7.5. This is because in methods compensating for the transport, a true pair has a small but non-zero residual due to detection noise. Consequently, the saturation threshold can be higher, allowing a legitimate pair with a large residual to still be associated. However, the cost of leaving two faces unpaired must remain above κ , otherwise it would be preferable to leave everything unpaired. In this way, a true pair with a large residual costs less than leaving its two faces unpaired. All costs cited in this thesis are based on methods that estimate the translation.

Seven pairing methods are compared in the rest of this work.

- **Order-based anchors.** Associates the i -th face on one side with the i -th on the other within each stacking level, without estimating any cross-view translation. After matching, this method applies the original system’s aggregation gate, which accepts a pair only if it passes three checks: torsion within 2.5° , nominal length within ± 10 mm of 6005 mm and tilt within 2.5° . However, since offline calibration introduces the bias on world X described in section 2.2, the length check rejects every pair and the method yields zero matches. To isolate the effect of that check, three variants are defined. The first removes the length check and is taken here as representative of the installed heuristic, since in the field the threshold is configured to 2000 mm instead of ± 10 mm. The second keeps the check but centres its tolerance on the estimated bias. The third applies none of the three checks and keeps only the order of the detections, which separates what the heuristic owes to its ordering from what it owes to its checks. All three variants rely only on the order produced by the detector, without estimating anything, which is why they are termed order-based anchors.
- **Free assignment.** Solves the assignment problem on the shared cost model without estimating any translation between the views Kuhn 1955 [5]. A constant penalty for unpaired faces is added to the cost matrix, allowing some instances to remain unpaired.
- **Aligned assignment.** Estimates a YZ translation between views for each scene via RANSAC consensus Fischler et al. 1981 [2], in which every pair of centroids proposes a translation and the one with the most inliers is selected. The assignment is then solved on the translated centroids and the translation is refitted once as the median residual of the matched pairs.
- **Monotonic alignment.** Introduces an order constraint. The method assumes that the relative order of the bars is preserved between the two views, so that a bar to the left of another in one view is also to its left in the other. Under the same translation estimate, both sides are sorted along Y and the assignment becomes a sequence alignment problem solved via dynamic programming Ohta et al. 1985 [10], with the order constraint imposed directly on the assignment. The constraint is enforced across the entire bundle rather than level by level, which makes it incorrect when a bar rests on top of others.
- **Hybrid method.** This is the approach proposed in this work. It applies monotonic alignment within each stacking level. The two views do not always report the same number of levels, for example when the top bar of a tall stack is recognised as a separate level in one view only, so in that case the levels are matched by their median height rather than by their index. A single ordered assignment over the entire bundle is used only when one of the views has a single level. It adds an explicit hypothesis for bar crossovers, which is accepted in two cases. In the first, both views independently observe the same number of bars raised relative to the rest of the bundle and those bars break the rigid motion of the bundle between the two views. In the second, a bar lying across the bundle rests on its neighbours at one end and on the bed at the other, so it is raised in one view only and has no counterpart near it once the translation between the views is applied. That bar is paired with the faces left unpaired on the other side using only orientation and distance, because the level and height terms of the cost model penalise precisely the elevation that identifies it.

2.5 The shift family and identifiability

The bars of a stacking level lie next to each other in a fixed order from left to right, with an almost constant pitch p within the bundle. The visible faces of a level are indexed along Y on each side, with the indices of side 2 offset so that the true correspondence pairs face i of side 1 with face i of side 2. For an integer k , the shift family is defined as

$$\pi_k = \{ (i, i + k) : \text{both indices exist within the level} \} \quad (2.2)$$

which associates face i of side 1 with face $i + k$ of side 2, leaving the $|k|$ faces at the boundary of each view unpaired. The true correspondence is therefore π_0 . Each π_k is a valid partial correspondence (section 2.4), so the shared cost model assigns a cost to each of them.

The shift family is important for three reasons:

Geometric degeneracy under exact regularity. When the pitch is constant, shifting the association by k positions and increasing the estimated translation by $k \cdot p$ leaves the YZ distances between paired faces unchanged. For a method that estimates a translation for each scene, the geometric terms of the cost therefore take the same value for π_k and for π_0 . This is the degenerate case of the unlabelled sensing problem Unnikrishnan et al. 2018 [15], in which the correspondence can be recovered only from generic measurements, so that a bundle with a constant pitch is the extreme case of non-generic measurements.

Unpaired detections and truncation. Since the bundle is finite, π_k leaves unpaired the $|k|$ faces at the boundary of each view, whereas π_0 pairs every bar that both views contain. When the two views contain the same bars, the true correspondence therefore leaves fewer detections unpaired than any shifted one, whatever the transport. Pitch irregularity adds a geometric difference between the two, because an irregular bundle shifted by a whole pitch no longer matches the other view. Truncation that differs between the two views has the opposite effect. When one view cuts off a bar that the other contains, the true correspondence also leaves detections unpaired, so that against a shift towards the truncated side only the geometric difference remains, which is small when the pitch is regular. Neither the regularity of the pitch nor the position of the bundle in the scan window is under the control of the vision system.

The family is a proper subset of possible errors. A repermutation (i.e., a permutation that alters the relative order) corresponds to a bar crossing over its neighbours. No π_k contains a correspondence of this type. On real scenes, and restricted to the detections that have an annotated partner, the only associations that the cost model ranks above the ground truth leave a crossing bar unpaired, so they also fall outside the family (section 4.4). On the full detection set, by contrast, the wrong outputs of the hybrid method are shifts that cost less than the annotated pairs (section 5.2).

2.6 The synthetic twin and its limits

The synthetic twin is a geometric simulator of the line. It generates bars by extruding the DXF cross-sections along the axis, arranges them into bundles with known 6D poses and simulates two virtual scanners by ray-casting onto the calibrated geometry. The output is saved in the same scene files used for real data, allowing the pipeline (loading, detection and pairing) to process them without modification. The bars are generated by the simulator, so their identity is known by construction and requires no annotation. This enables controlled

experiments that the real dataset cannot offer: variations in bar pitch regularity, inter-view transport, bundle position and the cost model margin relative to an artificially constructed member of the shift family. Figure 2.5 compares a real and a simulated scene.




Figure 2.5: Twin against real. The detector’s input for a real scene of the square profile and for its twin, a simulated scene with the same profile and the same number of bars, real above and simulated below. On the right, the real and twin medians for seven quantities. Each row uses its own scale.

The audit is conducted on the 56 scenes of the single category. Detection counts per scene-side agree at a Kolmogorov–Smirnov statistic of $D = 0.04$ and detection accuracy is comparable, with precision and recall of 0.974 and 0.964 on real scenes against 1.000 and 0.918 on the twin.

Point statistics, however, do not match, since the twin contains 46% of the real cloud size and only 19% of the real spread along world X (24 mm against 127 mm). Real bundles stagger the ends of their bars along the bar axis, whereas those of the twin are nearly flush.

Consequently, the twin is a stand-in for relative comparisons only and no conclusion that depends on the world-X channel is drawn from it. Every result obtained on the twin remains in the YZ projection and in the cost units of the model itself.

Chapter 3

Identifiability of the carriage geometry

The available offline calibration is valid in the YZ projection plane but incorrect along the X-axis, i.e., along the bar direction. Reconstructing the length of a bar by connecting the centroids of its two end faces yields a value systematically smaller than the nominal value of 6005 mm. The median residual is -1077.47 mm, with a very narrow IQR, $[-1080.36, -1076.02]$. The residual is therefore an almost constant offset rather than random noise. The default length tolerance of the original system’s pairing gate is ± 10 mm and none of the 425 pairs considered falls within it. This occurs because the carriage positions, which the controller reads at scan time, are not stored in the scan data available to this work. The offline calibration has a single knot per scanner, valid for one carriage position, so applying it evaluates the calibration at that position rather than at the actual one and introduces an error of about one metre along X.

While scene content alone recovers part of this geometry, it provably cannot recover the rest. The only metric datum that the recorded data preserves regarding a bar is its length, whose prior constrains only a single combination of the two carriage positions, leaving the orthogonal one free.

3.1 The length prior and its residual

Each bar is cut to a nominal length $L = 6005$ mm, so each pair of end faces, one per view, should reconstruct exactly that value.

Let s_1 and s_2 denote the positions of the two carriages along their respective translation axes and $w_1(s_1)$ and $w_2(s_2)$ the centroids in the world frame obtained by applying the calibration to those positions. The two views are also slightly offset in the projection plane. That offset is constant across the entire dataset along Z and varies scene by scene along Y (section 3.4), so the estimation involves a single unknown t_z and no global term along Y. For a pair k the residual is

$$r_k(s_1, s_2) = \left\| w_1^k(s_1) - \left(w_2^k(s_2) + (0, 0, t_z) \right) \right\| - L \quad (3.1)$$

The estimation of s_1 and s_2 minimises the median of the absolute value of the residuals:

$$(\hat{s}_1, \hat{s}_2) = \arg \min_{s_1, s_2} \text{median}_k |r_k(s_1, s_2)| \quad (3.2)$$

rather than the sum of squares because the pairs used for fitting include some uncertain associations whose large residuals could dominate a least-squares fit. The median absolute residual is far less sensitive to such outliers, making it the preferred choice here.

Since the objective is piecewise linear and exhibits non-differentiable points, the optimiser first performs a coarse grid search within the calibrated interval and then a local refinement, so that the result does not depend on an arbitrary initial point.

3.2 Identifiability of the carriage positions

The residual depends on the two carriage positions but not on each one independently. Owing to the structure of the model, if both carriages shift by the same amount along X , both reconstructed positions move together, so the distance between them stays the same. Every reconstructed bar length is therefore unchanged, as is every residual. A prior on the length therefore captures only the difference $s_1 - s_2$. The orthogonal direction, along which the two carriages slide together, belongs to the exact kernel of the Jacobian of the residual, so it is not a poorly observed direction that more data could improve. Since the observations carry no information about this direction, it cannot be recovered with more data, another optimiser or a different selection of scenes.

The degree to which the observed direction $s_1 - s_2$ is determined, by contrast, depends on the observations: on the number of pairs, their accuracy and their distribution. It is quantified by the curvature of the objective along that direction. The evaluation tool is the Gauss–Newton normal matrix $J^\top J$ of the residual with respect to the two parameters, divided by the number of pairs entering the fit to enable comparison across scenes of different sizes. It is the same observability analysis that the self-calibration literature applies to the parameters of a system Maye et al. 2016 [8], Nobre et al. 2019 [9], here reduced to a static problem with only two parameters.

Across all 425 pairs, the following eigenvalues are obtained

$$\lambda_1 = 2.008, \quad \lambda_2 = 2.91 \times 10^{-8}, \quad \lambda_1/\lambda_2 = 6.90 \times 10^7$$

The eigenvector associated with the largest eigenvalue, i.e., with the direction best determined by the data, is $(+0.708, -0.707)$, which corresponds to the direction $s_1 - s_2$ up to rounding errors. The matrix is obtained by linearising the residual around the estimate to which the alternation between calibration and assignment converges. Although linearisation is typically a local approximation whose result might depend on the evaluation point, this is not the case here because the null direction is exact across the entire parameter space and not only near the estimate. Consistently with this geometry, the second eigenvalue is zero within numerical precision.

Consequently, the objective is flat along that direction, so it has no interior minimum there and the search continues until it reaches the boundary of the calibrated domain. Whatever pair of individual positions an estimate reports depends on where that boundary lies and only the difference between the two is a measurement. Figure 3.1 shows this flat direction.

The same matrix also shows that no scene is geometrically more favourable than the others. The contribution of each pair to the curvature along the observed direction varies between 2.007 and 2.009 across all 108 scenes, i.e., it is practically constant. Each pair therefore carries almost the same information regarding the separation because the derivative of the distance between the two centroids with respect to a carriage position is close to 1 for any bar, regardless of its position, its inclination and the number of bars in the scene. The accuracy of the estimate depends on the number of pairs and their correctness, rather than the choice of favourable scenes.




Figure 3.1: Unobservable direction of the length prior. Contour lines of the quadratic form in the two carriage positions s_1 and s_2 , plotted starting from the two eigenvalues and the eigenvector reported by the observability analysis. The cost grows rapidly along the $s_1 - s_2$ direction (the observed one), while remaining flat along the orthogonal direction in which the two carriages slide together. With a curvature ratio of 6.90×10^7 , the curves extend without bound in the flat direction. The width of the plotted strip is an empirical dispersion level, 3.4 mm at real scale, magnified 22 times to render it visible. The dashed line marks the lowest knot of the calibration interval, i.e., the boundary at which a search along the flat direction stops.

3.3 Recovering the carriage separation

The estimation requires pairs as input, yet these derive from an assignment that itself depends on the calibration. Calibration and assignment are therefore estimated alternately:

1. assignment is performed, under the current calibration;
2. the positions $\hat{s}_1, \hat{s}_2$ and the residual offset are re-estimated on the obtained pairs;
3. the process is repeated until the assignment is identical between two consecutive iterations.

Here convergence means that the assignment has stopped changing, not that a numerical tolerance has been reached. Two iterations are needed.

As the unobserved direction predicts, $\hat{s}_2$ settles at -3339.0 mm, the lower knot of the calibration interval, rather than at an interior minimum. Since no data constrain that direction, this value only marks where the search stops, whereas the difference $\hat{s}_1 - \hat{s}_2$ is a genuine measurement.

Table 3.1: Bar length residual under the three calibrations. The residual is the reconstructed bar length minus the nominal 6005 mm. For the offline calibration, a one-shot fit on a fixed assignment and the joint self-calibration, the table reports the median residual, its interquartile range in millimetres and the number of pairs within the default ± 10 mm tolerance of the length check of the original system. The first and last rows use the pairs of the hybrid method, whereas the one-shot fit uses those of free assignment.

bar length residual	median	IQR	within ± 10 mm
offline calibration with single knot	-1077.47 mm	$[-1080.36, -1076.02]$	0 / 425
one-shot fit on fixed assignment	-1.48 mm	$[-5.11, +0.15]$	359 / 415
joint self-calibration	-0.15 mm	$[-2.99, +1.38]$	403 / 425

The first and last rows of table 3.1 use the hybrid method, on the offline calibration and on the recovered one, respectively. The middle row instead uses a free assignment and keeps the pairs fixed throughout the process. It therefore differs from the joint fit in two ways: the origin of the pairs and their ability to change. Its median is nominally closer to zero, but this is uninformative given that the median is precisely the quantity optimised by the objective. The informative columns are the dispersion and the fraction of pairs within ± 10 mm. With an IQR of 4.36 mm against 5.26 mm the joint fit is tighter, as expected, because it no longer fits the calibration to pairs that the calibration itself would have discarded.

The recovered separation is

$$\hat{s}_1 - \hat{s}_2 = 5403.3 \text{ mm}$$

and has an independent verification. Both the replay tool and the test harness of the original system use a default carriage position of 2701 mm per side, i.e., 5402 mm of separation, described in its documentation as a dashboard convention. The estimated value, obtained solely from scene content and the length prior without access to that constant, agrees with it within 1.3 mm.

None of this establishes that the recovered separation is accurate in an absolute sense. No independent metrology of this rig exists and the constant of the original system, however

close, is a software convention rather than a measured distance. The fit also has two free parameters and only one identifiable combination. A richer calibration model, with per-knot corrections or a rotational term, is not tested here and is not guaranteed to remain identifiable from the same prior.

Regarding uncertainty, the quantity to quote is the reproducibility across scenes rather than the repeatability across pairs. Estimating the separation independently on each scene yields a median of 5403.0 mm and an IQR of [5401.9, 5405.3]. A bootstrap on the joint fit that resamples bundles (section 2.1) gives [5402.8, 5403.8]. Figure 3.2 shows the separation estimated on each scene.




Figure 3.2: Cross-scene reproducibility of the carriage separation. The carriage separation estimated independently on each of the 108 scenes. The joint fit of 5403.3 mm is shown with the interquartile range of the per-scene values, along with the working constant of the original system (2701 mm per side), a value to which the fit never had access.

The dispersion across scenes indicates matching quality rather than calibration quality. The outliers comprise one scene with two pairs at one extreme (5378.2 mm) and one with four pairs at the other (5421.2 mm). A scene whose pairs are wrong will estimate whatever separation makes those pairs consistent, regardless of how many it contains.

3.4 Calibration in Z, transport in Y

A residual offset between views in the projection plane could stem from a miscalibrated axis or a transport of the bar between the two acquisitions. The former is constant, whereas the latter varies from scene to scene because the two views of a scene are acquired at different times while the conveyor moves. A variable offset therefore cannot be attributed to calibration.

The median cross-view offsets per scene, computed over 108 scenes for each axis, are reported below. Each scene's offset is the median across its pairs.

- Z: median +1.64 mm, IQR [+1.12, +2.30]. With an IQR width of 1.18 mm, the Z offset varies by little more than a millimetre across scenes, a magnitude compatible with a constant calibration offset.
- Y: median -5.36 mm, IQR [-47.56, +21.04]. The IQR width is 68.60 mm, nearly sixty times that of Z, while 62 of the 108 scenes exhibit an offset exceeding ± 30 mm. The dispersion is thus wider and heavy-tailed, with the variability concentrated along the transport axis.

Consequently, the Z component is attributed to calibration, while the Y component is attributed to transport. This is why the residual described in section 3.1 includes a dataset-wide constant term along the Z-axis, but none along the Y-axis. This distinction forms the basis of a choice already adopted by the pairing methods: aligned assignment, monotonic alignment and the hybrid method all estimate a scene-specific YZ translation via RANSAC consensus, rather than relying on calibration along the Y-axis.

3.5 Metric recovery and the combinatorial ambiguity

Recalibration does not alter the relative performance of the pairing methods. The counting F1 varies by at most 0.002 for each of the methods re-scored. The exception is the order-based anchor with the length check of the original system, which jumps from 0.000 to 0.917. This jump comes from removing the -1077 mm bias, which made every pair fail the check. Once the bias is removed the check can be passed again, although this does not mean that it separates correct pairs from incorrect ones (section 4.1).

None of the results presented in this chapter addresses the ambiguity that forms the core subject of this thesis. Recovering the separation solves the metric problem of reconstructing a cross-view pair with the correct length. The combinatorial problem remains unchanged, as a bar reconstructed with the correct length provides no indication of which among several identical neighbouring elements was actually paired with it. This is the first of the properties of the shift family stated in section 2.5: under exact regularity, geometry does not distinguish a shifted correspondence from the true one. The metric recovery in this chapter does not affect that property because it operates on the calibration rather than the combinatorial ambiguity.

Chapter 4

Shift ambiguity after self-calibration

Self-calibration removes the metric error of the rig, so that a correct cross-view pair reconstructs the nominal bar length (section 3.3). The length of a pair, however, gives no indication of which of several identical neighbouring bars a face belongs to. This combinatorial ambiguity is described by the shift family π_k of section 2.5, which pairs each face on one side with the face k positions further along on the other.

4.1 Sensitivity of the length check to a one-bar shift

The aggregation gate of the original system accepts a candidate pair only when the distance between the centroids of its two end faces lies within ± 10 mm of the nominal bar length L (section 2.4). Of the three checks in this gate, it is the only one based on a property of the product itself, whereas the torsion and tilt checks, like every term of the shared cost model, compare the geometry of the two faces.

Consider a member π_k of the shift family, which pairs the end of bar i on one side with the end of bar $i + k$ on the other. The reconstructed segment then joins two different bars and becomes the hypotenuse of a right-angled triangle whose legs are the bar length L and the lateral offset kp , where p is the bar pitch. The reconstructed length therefore changes by

$$\Delta\ell(k) = \sqrt{L^2 + (kp)^2} - L \approx \frac{(kp)^2}{2L} \quad (4.1)$$

with $L = 6005$ mm, so the change grows with the square of the offset. Table 4.1 gives $\Delta\ell$ for a one-bar shift at the measured pitch of each family (fig. 2.4), in millimetres and in units of $\sigma = 3.23$ mm, the standard deviation that a normal distribution would need to reproduce the interquartile range of the length residual after self-calibration (section 3.3).

Since the pitch is much smaller than the bar length, a one-bar shift changes the reconstructed length by at most about 1 mm, against a tolerance of ± 10 mm.

Setting $\Delta\ell = 10$ mm in the formula gives the lateral offset at which the check would react, $kp \approx 347$ mm, which is about 9 pitches for the square profile, 5 for the L-profile and 3 for the complex profile, the family with the largest pitch. More than nine clouds in ten have a narrower scan window than this (median 258 mm, 90th percentile 337 mm), so only

Table 4.1: Length change caused by a one-bar shift. For each family, the measured bar pitch, the change $\Delta\ell$ in the reconstructed length when a pair joins the ends of two adjacent bars and the same change in units of $\sigma = 3.23$ mm, the standard deviation of a normal distribution with the same interquartile range as the length residual after self-calibration. The length check of the original system has a tolerance of ± 10 mm.

family	pitch (mm)	$\Delta\ell$, one pitch (mm)	$\Delta\ell/\sigma$
L-profile, second acquisition	38.6	0.12	0.04
square profile	40.2	0.13	0.04
L-profile	69.2	0.40	0.12
heavy profile	96.0	0.77	0.24
complex profile	118.4	1.17	0.36

the widest ones, up to 414 mm (section 2.2), could contain such an offset. A threshold of $3\sigma \approx 9.7$ mm in place of the fixed ± 10 mm gives almost the same offset (341 mm), so the limit lies in the geometry of the check and not in the choice of tolerance.

4.2 The ambiguity as a function of cross-view transport

The ambiguity might be expected to depend on the cross-view transport and in particular to become more severe when the transport approaches a whole multiple of the pitch. Correspondence recall is not a suitable quantity for testing this, because it mixes three effects: whether the detector found the bars, whether the two views contain the same set of bars and whether the pairing objective can tell the true correspondence from a shifted hypothesis. Only the third effect is the ambiguity, so the experiment measures the objective directly, on synthetic scenes in which bar identity is known by construction (section 2.6). The synthetic twin reproduces the input of the detector but not the point cloud behind it, so the results of this section and of section 4.3 hold in the YZ projection and in the cost units of the model.

The objective is the one that monotonic alignment minimises and that the hybrid method solves within each stacking level. It scores a complete association, called a hypothesis here, as the sum of two terms. The geometric term sums the pair costs $c(i, j)$ of section 2.4 over the matched pairs, after aligning the two views by the median cross-view translation of those same pairs, so that each hypothesis is scored at the translation that suits it best. The counting term adds 3.75 for each detection that the hypothesis leaves unpaired on either side, half of the cost of 7.5 for leaving both faces of a pair unpaired (section 2.4). The experiment scores both the true correspondence π_0 and the one-bar shift π_1 in this way and reports

$$\text{margin} = \text{cost}(\pi_1) - \text{cost}(\pi_0), \quad (4.2)$$

which is positive when the true correspondence has the lower cost.

The sweep uses a bundle of five square-section bars. Its pitch is among the finest in the dataset and its 90° symmetry period leaves the orientation term with the least information, which makes it the hardest case for the objective. The cross-view transport goes from 0 to 3 whole pitches in steps of a quarter pitch and is combined with two placements of the

bundle and three levels of pitch jitter, 0, 2 and 10 mm. The jitter is the standard deviation of a Gaussian perturbation added to the lateral position of each bar in an otherwise regular bundle. In the centred placement no bar leaves the scan window at any transport, which keeps truncation out of the test, while in the other placement the same bundle sits at the edge of the window.

The margin against π_1 shows no periodicity. At both placements and at every jitter level, the median margin at whole multiples of the pitch differs from the median between them by at most 0.01 cost units, against margins between 7.5 and 8.6. Under regular spacing the ambiguity is therefore present at every transport, which is the degeneracy that section 2.5 states as the first property of the shift family. Figure 4.1 shows the centred sweep.

The centred sweep shows a single change, at the end of the range. At the largest transport the margin narrows slightly, from 7.69 to 7.48 with 2 mm of jitter and from 8.60 to 8.17 with 10 mm. No bar leaves the window and the number of unpaired detections stays constant, so the change comes from the geometric terms of the two hypotheses. At three pitches the true pairs fit slightly worse than at smaller transports. Without jitter the shifted pairs fit worse by the same amount and the margin does not change, whereas with jitter they worsen less or even improve, so the true correspondence loses part of its advantage. The narrowing does not appear at one or two pitches, where a periodic effect would place it.




Figure 4.1: The ambiguity sweep. Cost margin between a one-bar shift and the true correspondence as a function of the cross-view transport, for a bundle of five square-section bars centred in the scan window, so that no bar ever leaves it, at three levels of pitch jitter. Each point is the median over three random seeds. The dotted lines mark whole multiples of the pitch, where a periodic ambiguity would show a dip, whereas the margin stays flat across all of them and drops only at the end of the range for the two jittered bundles.

The two terms of the objective explain why the margin stays constant. Their median values over the centred sweep are given in table 4.2.

At zero jitter geometry contributes almost nothing to the margin. A shifted hypothesis estimates its own translation, which absorbs the transport completely, so under perfectly regular spacing its pairs fit as well as the true ones. The whole margin comes from the

Table 4.2: Geometric and counting terms of the objective. Median values in cost units for the true correspondence π_0 and the one-bar shift π_1 , with the median margin cost $\text{cost}(\pi_1) - \text{cost}(\pi_0)$, for each level of pitch jitter. The medians are taken over the centred sweep, in which a synthetic bundle of five square-section bars stays inside the scan window while the cross-view transport goes from 0 to 3 pitches.

jitter	geometric, π_0	geometric, π_1	counting, π_0	counting, π_1	margin
0 mm	0.04	0.06	0.00	7.50	7.52
2 mm	0.07	0.22	0.00	7.50	7.69
10 mm	0.05	1.12	0.00	7.50	8.59

counting term because in a bundle of n bars a one-bar shift pairs only $n - 1$ of them and the two detections it leaves unpaired cost 3.75 each.

Pitch jitter makes geometry informative. At 10 mm the geometric term of the shifted hypothesis rises to 1.12 while that of the true correspondence stays at 0.05, because an irregular bundle cannot be slid by one pitch and still fit. Irregularity therefore plays the role that genericity of the measurements plays in unlabelled sensing Unnikrishnan et al. 2018 [15], since a bundle with constant pitch is the degenerate case and jitter restores the discriminating power of geometry. This is the mechanism that section 2.5 states as the second property of the shift family.

The sweep covers a single family, a single bar count and three seeds per cell. This is enough to establish whether the mechanism is periodic but not to estimate its size for other families. The margin is also a cost difference in the units of the model and not an error rate, so a wide margin can still be lost to a detection failure upstream.

4.3 Truncation and transport

Since the margin comes from the counting term, anything that changes the number of unpaired detections changes the margin, including a scan window that cuts off a bar. Truncation is common on this rig, since 30 of the 108 scenes contain at least one face without a partner (section 2.1) and six L-profile bars at the measured pitch already span about 415 mm, against a median scan window of 258 mm (section 2.2).

Two matched synthetic suites separate truncation from transport. Both use a bundle of four L-profile bars and sweep the transport from 0 to 1.5 pitches, pooling six jitter levels between 0 and 20 mm. In the first suite the bundle sits at the edge of the scan window, so that transport carries bars in and out of view, while in the second the same bundle is centred with clearance on both sides, so that no bar leaves the window. For each transport, table 4.3 reports the correspondence recall of the hybrid method over the bars detected on both sides and the number of scene-sides, out of 36, on which the number of detections differs from the number of faces present. This count measures how often truncation changed what a view sees.

The centred bundle keeps perfect recall at every transport, so within this bundle size the magnitude of the transport causes no damage on its own and the loss of recall at the edge comes from truncation.

The edge column adds a more specific result. At zero transport it has the highest count mismatch in table 4.3 and still perfect recall, because both views truncate the same bar,

Table 4.3: Truncation against transport. Synthetic bundle of four L-profile bars placed at the edge of the scan window or centred in it, with the cross-view transport expressed as a fraction of the bar pitch. For each transport, the correspondence recall of the hybrid method over the bars detected on both sides and the number of scene-sides, out of 36, on which the number of detections differs from the number of faces present. Each row pools six jitter levels between 0 and 20 mm.

transport / pitch	edge recall	edge count mismatch	centre recall	centre count mismatch
0	1.000	35/36	1.000	0/36
0.25	1.000	20/36	1.000	0/36
0.50	1.000	18/36	1.000	0/36
0.75	0.889	18/36	1.000	0/36
1.00	0.944	18/36	1.000	0/36
1.25	0.944	18/36	1.000	0/36
1.50	0.944	18/36	1.000	0/36

their visible sets coincide and no wrong pairing is possible. A non-zero transport makes the truncation differ between the two views and recall follows this asymmetry rather than the distance travelled, since it is lowest at an intermediate transport of 0.75 pitches and not at the largest one. Figure 1.2 shows exactly this configuration. As the decomposition of section 4.2 explains, when the two views truncate different bars the true correspondence itself leaves detections unpaired and pays the counting term, so it loses the advantage that this term normally gives it over a shifted hypothesis.

The sweep of section 4.2 measures this effect directly on its bundle of five square-section bars. At the edge placement, once the transport reaches three quarters of a pitch the two views detect different numbers of bars and the margin against the shift towards the truncated side, π_{-1} , collapses. When the two views detect the same number of bars its median is 7.3 to 8.3, almost all from the counting term. When they do not, both hypotheses leave the same number of detections unpaired and the median margin is only the geometric difference, 0.02 cost units for a regular bundle, 0.19 with 2 mm of jitter and 1.11 with 10 mm. On 5 of the 117 edge scenes the shift costs less than the true correspondence. The margin against the shift in the other direction stays between 7.5 and 8.3 throughout.

4.4 Cost of the true correspondence on real scenes

The synthetic results describe what the objective can distinguish under controlled conditions. On real data the corresponding question is whether the annotated correspondence of a scene costs less than every alternative association of the same detections.

The alternatives come from two sources. The first is a search in which every pair of detections across the two sides proposes a cross-view translation, the order-constrained assignment is solved under each proposal, each solution is scored at its own median translation as in section 4.2 and duplicates are removed. On small regular bundles this search does not reach the shift family, because under every proposed translation the solver returns the true correspondence. The second source therefore adds the members $\pi_{\pm 1}$ and $\pi_{\pm 2}$ of the shift family explicitly, built within each stacking level. Because these shifts are constructed rather than found, the comparison tests the objective itself and not a particular search, since a constructed shift that cost less than the true correspondence would reveal a fault of the objective.

The annotated correspondence is scored with the same objective and the same constants, with all hypotheses restricted to the detections covered by the annotation, so that a bar the annotator did not claim cannot add to the cost of the true correspondence. For each scene the quantity of interest is the difference between the cost of the cheapest alternative and that of the true correspondence.

The analysis covers the 102 scenes with at least two resolved claims. Its main result is that in every one of them the true correspondence costs less than a one-bar shift. The median margin against $\pi_{\pm 1}$ is about 7.3, close to the 7.5 of the counting term in the synthetic sweep, so the real scenes behave as section 4.2 describes, with the margin coming mainly from the two detections that a shift leaves unpaired. Because every hypothesis is restricted to the annotated detections, a constructed shift always leaves two of them unpaired. The comparison therefore shows that on real scenes the geometric terms are never large enough to outweigh the counting term. It leaves out the detections without an annotated partner, which is where truncation acts, even though every scene on which the hybrid method fails completely contains at least one of them (section 5.2).

In two scenes, however, an alternative costs less than the true correspondence, by less than half a cost unit. Both scenes belong to the crossed category of the L-profile and are two acquisitions of the same bundle. In both, the cheapest alternative is the true correspondence with one pair removed, which keeps every other true pair and leaves unpaired the two ends of a bar lying across the bundle. This bar rests on the bundle at one end and on the bed at the other, so its face sits about 30 mm above the others in one view and level with them in the other. Its true pair therefore collects the stacking-level penalty, the height-outlier penalty and a distance cost close to its cap, so the total exceeds the cost of leaving both ends unpaired. In none of the 102 scenes does an alternative that pairs a bar with the wrong partner cost less than the true correspondence. The two exceptions are therefore a missed pair on a crossing bar, the kind of error outside the shift family described by its third property (section 2.5), which the crossing hypothesis of the hybrid method handles (section 2.4).

The comparison with section 4.1 suggests one further question: whether scenes with a larger implied length change also have a larger margin. Over the scenes whose cheapest alternative belongs to the shift family, the rank correlation between $\Delta\ell$ and the margin is weakly positive (Spearman's $\rho = 0.25$). This does not show that the length cue is at work, since the objective has no length term. Both quantities depend on the pitch because $\Delta\ell$ grows with it by construction and a wider pitch also separates the bars in the distance term that the objective does contain. The correlation therefore reflects this common factor rather than an effect of length. The metric cue around which the original system was designed is therefore absent from the objective that decides the association, while section 4.1 shows that it could not resolve the error even if it were present.

Chapter 5

Evaluation of the pairing methods

The pairing methods of section 2.4 are evaluated on the 108 real scenes with the metrics of section 2.1.

All results refer to the candidate-pool detector (section 2.3), except where sections 5.3 and 5.4 use the two-pass suppression detector. The order-based anchor with the original length check is excluded from all comparisons. Under the offline calibration, the X-axis bias of about -1077 mm (section 3.3) makes the ± 10 mm length check fail on every candidate pair so the method yields no pairs. After self-calibration its counting F1 returns to 0.917 (section 3.5), which shows that the missing pairs come from the calibration error rather than from the method itself. In its place, the comparisons use the three variants defined in section 2.4: one without the length check, one with the check re-centred on the estimated error (the debiased check) and one that keeps the order of the detections without applying any check.

5.1 Counting and correspondence metrics on the same detections

The simplest way to evaluate a pairing stage is to compare the number of pairs produced by a method with the number expected from the count annotation. This metric needs no correspondence annotation but cannot tell a correct pair from a swapped or shifted one (section 2.1). Table 5.1 reports both metrics for each method on the same detections across the 108 scenes, along with the inversion count at the same stacking level and the IPAA-1.0 (i.e., the fraction of scenes with entirely correct pairing).

The counting metric discriminates poorly between methods, since thirteen of the twenty-one pairs of methods have overlapping 95% intervals on counting F1. It separates the two order-based anchors that retain a check from all other methods, with the exception of monotonic alignment, but fails to rank the remaining five methods among themselves.

When the counting metric does distinguish two methods, correspondence recall does not always confirm the ranking. While it confirms the ordering for the hybrid method and the check-free anchor, the two anchors with checks yield a lower counting F1 (0.917 and 0.933 against 0.973) yet achieve higher correspondence recall (0.816 and 0.835 against 0.762 and 0.779) compared to free and aligned assignment. A paired bootstrap on the bundles reveals no significant differences between these recall values. The correspondence annotation therefore does not show that the anchors pair better than free or aligned assignment. It shows instead that the gap in counting F1 does not reflect a real difference in pairing accuracy.

Table 5.1: Counting and correspondence metrics of the pairing methods. All methods pair the detections of the candidate-pool detector on the 108 real scenes. Counting precision (P), recall (R) and F1 compare the number of pairs with the count annotation and the 95% interval of counting F1 comes from a bootstrap over bundles. Inversions count side-by-side pairs at the same stacking level whose Y-order is inverted between the two views, correspondence recall is the fraction of resolved claims reproduced with the correct bar identity and IPAA-1.0 is the fraction of scenes paired entirely correctly.

method	counting P	counting R	counting F1	F1 95% CI	inversions	correspondence recall	IPAA-1.0
order-based anchor, debiased check	1.000	0.847	0.917	[0.891, 0.941]	0	0.816	0.731
order-based anchor, no length check	0.997	0.876	0.933	[0.909, 0.954]	0	0.835	0.741
order-based anchor, no checks	0.995	0.948	0.971	[0.958, 0.984]	0	0.872	0.843
free assignment	0.995	0.953	0.973	[0.961, 0.985]	59	0.762	0.676
aligned assignment	0.995	0.953	0.973	[0.961, 0.985]	40	0.779	0.741
monotonic alignment	0.998	0.928	0.961	[0.940, 0.979]	0	0.865	0.796
hybrid method	0.995	0.955	0.975	[0.961, 0.987]	0	0.929	0.889

Free assignment shows why, since it reaches almost the highest counting F1 with the lowest correspondence recall. It pairs faces without any order constraint, so 59 of its pairs are same-level inversions that give the expected number of pairs with the wrong identities.

This is an instance of the discrepancy recently described for multi-view association Shelukhan et al. 2026 [13], where a better score on a proxy metric does not necessarily mean a better assignment.

5.2 Structure of the pairing errors

If a method made independent errors on individual pairs, most scenes with an error would be only partly wrong. To test this hypothesis, each scene is classified as clean when all resolved claims are exact, wrong when none are and partial in all other cases (table 5.2). A wrong scene is considered a pure shift when all its claims are shifted by the same non-zero index offset k within their respective stacking level. In that case the output of the method for that scene corresponds to element π_k of the shift family defined in section 2.5, rather than an arbitrary permutation.

In the case of the hybrid method, all five wrong scenes contain at least two claims and are pure shifts by exactly one bar. Together they account for 19 of the 29 wrong claims, which means that most of the method’s error comes from whole scenes shifted at once rather than from noise on individual pairs. A pair-based metric describes this error process poorly, since the failing unit is the scene itself. For this reason, every confidence interval in this work involves resampling whole scenes grouped into bundles (section 2.1).

These five wrong scenes share a common property. Of the 102 scenes with at least two resolved claims, 59 have an annotated partner for every detection on both views and the hybrid method makes no error on any of them, whereas the five wrong scenes are all among the remaining 43, where at least one detection lacks an annotated partner. Such a detection

Table 5.2: Classification of the scenes by pairing outcome. Each of the 108 real scenes is clean when all resolved claims are correctly reproduced, wrong when none is and partial otherwise, with the detections of the candidate-pool detector. The wrong scenes are also counted when they have at least two resolved claims and when they are pure shifts, that is, when all claims are shifted by the same non-zero offset k within their stacking level. The last column breaks the pure shifts down by k .

method	clean	partial	wrong	wrong, ≥ 2 claims	pure shift	shift sizes
order-based anchor, no length check	80	12	16	14	12	$k = -1$: 5, $k = +1$: 7
order-based anchor, no checks	91	3	14	13	10	$k = -1$: 3, $k = +1$: 7
free assignment	73	19	16	16	3	$k = -1$: 1, $k = +1$: 2
aligned assignment	80	14	14	14	5	$k = -1$: 3, $k = +1$: 2
monotonic alignment	86	16	6	6	5	$k = -1$: 3, $k = +1$: 2
hybrid method	96	7	5	5	5	$k = -1$: 3, $k = +1$: 2

may be a bar that the other view truncates, a bar whose face the detector misses in the other view, a false detection or a bar that the annotation does not claim. In many scenes the notes of the annotation record which case applies. In the five wrong scenes, inspection of the point clouds shows that every detection without a partner is a bar that the other view truncates.

To explain the failure in these scenes, the output of the method and the annotated correspondence are both scored with the objective of section 4.2 on every detection of the scene, leaving unpaired each detection that the annotation does not pair. In all five scenes the wrong output costs less than the annotated correspondence so the failure comes from what the objective can distinguish and not from the search that minimises it.

The five scenes split into two groups. In the first group, consisting of two scenes, each view contains at least one bar that the other does not see. The annotation thus leaves unpaired detections on both sides, whereas the shift associates more of them and costs 6.6 to 7.4 cost units less than the annotated correspondence. In the second group, consisting of the other three scenes, the shift associates the same number of detections as the ground truth. Here the counting term is the same for both associations so the difference comes from geometry alone, which favours the shift by only 0.02 to 0.60 cost units. The same behaviour occurs in synthetic scenes at the edge of the scan window (section 4.3). In those scenes truncation cuts a bar in one view but not in the other, so even the true correspondence leaves an unpaired detection, the count advantage vanishes and the margin depends solely on geometry.

Consequently, these errors cannot be eliminated with a better search because in these scenes the objective already favours the wrong association. The analysis, however, identifies a clear risk condition, since the hybrid method fails completely on a scene only when at least one detection lacks an annotated partner.

5.3 Correspondence accuracy of the pairing methods

To determine whether differences in correspondence recall exceed sampling variation, each method is compared with the hybrid method by a paired bootstrap over the 96 bundles formed by the 108 scenes (section 2.1). Each of the 10,000 replicates draws a set of bundles with replacement and evaluates all methods on that identical set. Consequently, variation due to scene selection cancels out when taking the difference between two methods.

Table 5.3: Correspondence recall compared with the hybrid method. Correspondence recall of each pairing method on the 108 real scenes with the candidate-pool detector and its difference Δ from the hybrid method, whose recall is 0.929. The 95% interval and p come from a paired bootstrap with 10,000 replicates over the 96 bundles, in which every replicate evaluates all methods on the same resampled bundles.

method	correspondence recall	Δ vs hybrid	95% paired CI	p
order-based anchor, no checks	0.872	-0.057	$[-0.133, +0.016]$	0.13
monotonic alignment	0.865	-0.064	$[-0.127, -0.018]$	<0.001
order-based anchor, no length check	0.835	-0.093	$[-0.172, -0.016]$	0.02
order-based anchor, debiased check	0.816	-0.113	$[-0.192, -0.034]$	0.007
aligned assignment	0.779	-0.150	$[-0.259, -0.065]$	<0.001
free assignment	0.762	-0.167	$[-0.270, -0.067]$	0.001

The hybrid method achieves the highest correspondence recall (table 5.3) and differs significantly from five of the six competing methods. Since this conclusion rests on six tests at once, the comparisons are corrected with the Holm–Bonferroni procedure Holm 1979 [4], which keeps the probability of even one false separation at 5% or less and remains valid when the tests are dependent, as they are here because they share the same bootstrap replicates. All five comparisons whose confidence intervals exclude zero survive this correction. The same pattern holds for IPAA-1.0, where the gap compared to monotonic alignment is -0.093 with a 95% interval of $[-0.153, -0.043]$.

The sole exception is the order-based anchor without checks, which associates detections within each stacking level based strictly on order. Its recall is 0.057 lower but the confidence interval for the difference includes zero. Similarly, on IPAA-1.0, the gap (-0.046 with interval $[-0.123, +0.028]$) is not statistically significant. Switching to the two-pass suppression detector gives the same picture, since the hybrid method again leads in recall (0.911) and differs significantly from all methods except this anchor (0.856, difference of -0.055 , interval $[-0.124, +0.003]$). Given this sample size, ordering alone serves as a baseline that none of the more elaborate methods significantly outperform. However, the order-based anchor without checks fails completely on more than twice as many scenes as the hybrid method (13 versus 5 among those with at least two claims, section 5.2).

The gap between the two methods comes from a small and specific set of scenes. On the 64

scenes in which both views report the same number of detections and the crossing hypothesis (section 2.4) is not accepted, the two methods associate the same number of claims correctly. Where the numbers of detections differ, the hybrid method does better on 7 scenes and worse on 6 and gains 17 claims while losing 18, which gives neither method a real advantage. The hybrid method gains 24 claims on the 6 scenes in which the crossing hypothesis is accepted and the order-based anchor without checks associates no claim correctly. These scenes therefore account for the whole difference of 23 claims (378 against 355) and are also among the 13 on which the anchor fails completely.

The hybrid method applies monotonic alignment within each stacking level while incorporating the crossing hypothesis (section 2.4). Its advantage over monotonic alignment stems from structural configurations that a single ordered association across the entire bundle cannot represent, namely bundles with several stacking levels, including those whose top bar is recognised as a separate level in one view only, and bars lying across the bundle, whether raised in both views or at one end only. The hybrid method never scores lower than monotonic alignment on any scene and scores higher only on scenes of these kinds.

The crossing hypothesis is accepted in 9 of the 108 real scenes. Two synthetic suites isolate this mechanism under controlled conditions. In the first suite, the crossing bar is raised along its entire length and the hybrid method correctly associates every bar in every cell. Free and aligned assignment, which lack the crossing hypothesis, achieve the same result, so this suite confirms that the hypothesis triggers when expected without showing a gain in accuracy. In the second suite, one end of the bar rests on the bundle while the other extends onto the bed beyond the bundle, as in the real scenes discussed in section 4.4. Here, the hybrid method again correctly associates every bar in every cell, whereas every competing method fails in at least half of the cells.

5.4 Stability under a change of detector

The comparison of section 5.3 uses a single detector (the candidate-pool detector) at its standard operating thresholds and does not show whether the methods remain robust when the detector changes. The test is possible because claims are anchored to YZ points rather than to detection indices (section 2.1) so the same ground truth scores both detectors. Holding the ground truth and scenes fixed, the candidate-pool detector is replaced with the two-pass suppression detector (section 2.3) to measure the resulting change in recall for each method.

The comparison uses the 108 scenes in which both detectors resolve at least one claim, evaluating confirmed detections exclusively. In table 5.4, Δ represents the correspondence recall under the candidate-pool detector minus that under the two-pass suppression detector. The final column compares each method against the order-based anchor without a length check, the variant selected in section 2.4 to represent the installed heuristic. This column reports $|\Delta(\text{anchor})| - |\Delta(\text{method})|$, which yields a positive value when a method exhibits greater stability (i.e., smaller change) than the anchor. Both confidence intervals are derived from a paired bootstrap over the bundles. Figure 5.1 shows the recall of each method under the two detectors.

Only the two order-based anchors that keep the checks of the original system change substantially, both losing roughly 0.50 in recall, which is more than half of their performance. In contrast, all other methods change by less than 0.10, with the 95% interval for Δ containing zero for every method except free assignment. The anchor without checks belongs to this second group even though it pairs the detections in exactly the same order as the




Figure 5.1: Effect of a change of detector. Correspondence recall for each method across the two detectors using identical ground truth and scene sets. Colours distinguish four methodological groups: methods using the shared cost model with a per-scene translation estimate, free assignment (which employs the cost model without a translation estimate), the two order-based anchors retaining acceptance checks from the original system and the check-free order-based anchor. Dashed lines with open markers serve solely to resolve visually overlapping trajectories. Only the two anchors with acceptance checks lose most of their recall, whereas the order-based anchor without checks remains as stable as the cost-model methods.

Table 5.4: Stability of correspondence recall under a change of detector. Δ is the correspondence recall with the candidate-pool detector minus that with the two-pass suppression detector, on the 108 scenes in which both detectors resolve at least one claim and with the same ground truth. The difference from the anchor is $|\Delta(\text{anchor})| - |\Delta(\text{method})|$, where the anchor is the order-based anchor without a length check, so a positive value means that the method changes less than the anchor. Both 95% intervals come from a paired bootstrap over bundles.

method	Δ	95% paired CI	difference from the anchor	95% paired CI
monotonic alignment	+0.004	[−0.058, +0.065]	+0.515	[+0.406, +0.577]
order-based anchor, no checks	+0.016	[−0.021, +0.053]	+0.502	[+0.418, +0.579]
hybrid method	+0.018	[−0.049, +0.082]	+0.501	[+0.397, +0.576]
aligned assignment	−0.022	[−0.108, +0.063]	+0.496	[+0.374, +0.574]
free assignment	+0.068	[+0.017, +0.122]	+0.451	[+0.353, +0.553]
order-based anchor, debiased check	+0.507	[+0.431, +0.583]	+0.012	[−0.011, +0.037]
order-based anchor, no length check	+0.519	[+0.439, +0.597]	reference	

anchor without the length check. The two differ only in the torsion and tilt checks, which compare the matching angle and the face normal of the two ends of a candidate pair. The loss of the installed heuristic under a change of detector therefore comes from these checks, whose inputs depend on how each detector estimates angles and normals, rather than from order-based pairing itself. The performance difference between the two anchors (+0.502, CI [+0.418, +0.579]) survives the Holm–Bonferroni correction alongside those of the four cost-model methods.

One caveat concerns the size of the effect but not its direction. Since the anchors were placed on the output of the candidate-pool detector, claims are re-linked to the two-pass suppression detector with a median error of 0.96 mm and fewer of them are resolved (section 2.1). This affects every method in the same way, which is why the comparison measures how much each method changes rather than which detector is more accurate. Moreover, the re-linking error cannot plausibly explain a recall drop of 0.50 that appears only in the methods with acceptance checks.

5.5 Comparison with the original system's pairing criterion

To contextualise these findings, the same 108 scenes were reprocessed using the replay tool of the original system, comparing its association decisions against the correspondence annotations. Reprocessing was conducted using offline calibration, the default carriage position of the original system (2701 mm per side, section 3.3) and a length tolerance of 2000 mm, the value configured on a machine in service. Across 407 claims in the 108 scenes, the replay tool achieves a correspondence recall of 0.531 and a precision of 0.818. This recall is below that of every method evaluated in this chapter, since even free assignment, the weakest of them, reaches 0.762 on the same claims.

Although these parameters were chosen in this work, they do not affect face detection, which takes place in the YZ projection, where the tool finds almost all faces (precision 0.996, recall 0.898). Correspondence recall is therefore lost in the pairing decision, which compares the 3D distance between centroids with the expected bar length. Since this distance depends mainly on the X coordinate, along which the offline calibration introduces an error (section 2.2), this error is probably the main reason for the lower recall.

Chapter 6

Conclusion

This thesis addresses a problem that arises when two 3D scanners observe opposite ends of a bundle of long, identical metal bars. Because the two scanners share no field of view and the bars are identical and almost evenly spaced, a single end face does not identify its bar. However, a robot tasked with grasping each bar at both ends must know which end in one view corresponds to which end in the other. The goal of this work is thus to establish which part of this correspondence the data can determine. This distinction is also important for evaluation, as a single accuracy metric cannot distinguish between errors that can be eliminated with an improved method and errors that are fundamentally unavoidable given the spatial configuration.

The approach establishes what the data can determine, recovers this part and measures the shift ambiguity that remains. This approach was applied to rig calibration (chapter 3), bar association (chapter 4) and the evaluation of pairing methods (chapter 5). For the rig, it is possible to recover the separation between the carriages but not their absolute position. For association, it is possible to recover the identity of the bars when both views contain the same bars but not always when the scan window truncates the two views differently. Association errors therefore depend on specific scene conditions and are not uniformly distributed across all scenes.

6.1 Summary of the results

Calibration of the rig. The nominal length of the bars determines the separation between the two scanner carriages but not their absolute position, since moving both carriages together leaves the reconstructed length of each bar unchanged. Recovering the separation from the data eliminates an error of about one metre along the bar axis and reduces the median length residual from -1077.47 mm to -0.15 mm. The recovered separation agrees within 1.3 mm with a constant defined independently in the original system.

The shift ambiguity. Once the rig is calibrated, the only remaining ambiguity is a shift of the entire association by an integer number of bars (section 2.5). When the bars are evenly spaced, a shifted association has practically the same geometric cost as the true one, so the two differ only in the number of detections left unpaired. This count favours the true association when both views contain the same bars, whereas when the scan window truncates the bundle differently across views even the true association leaves detections unpaired and the count no longer separates them. Furthermore, the length check of the original system cannot detect this shift, since pairing the ends of two adjacent bars changes the reconstructed

length by at most about one millimetre, against a tolerance of ± 10 mm.

Evaluation on real data. Across the 108 real scenes, the hybrid method achieves the highest correspondence recall (0.929). The method commits no errors in the 59 scenes where every detection has a partner, so its errors occur only in scenes containing detections without an annotated partner. In particular, the five scenes on which it fails completely are all shifted as a whole, contain at least one detection without an annotated partner and give an incorrect association with a lower cost than the annotated one. Because errors of this type, like swaps, can preserve the expected number of pairs (section 2.1), the counting metric, which only tracks how many pairs a method produces, fails to detect them and cannot rank the methods reliably.

6.2 Implications for pairing and evaluation

Relation to previous work. To the best of our knowledge, no previous work addresses this scenario. Methods for cross-view association typically rely on overlapping views, where geometric constraints connect the same surface in both images Sun et al. 2025 [14]. Alternatively, they rely on objects whose appearance varies from one instance to another Cai et al. 2020 [1], Schraml et al. 2015 [12]. These assumptions do not hold for identical bars observed from opposite ends. Identifiability, usually studied in the context of calibration Maye et al. 2016 [8], Nobre et al. 2019 [9], is applied in this thesis to two problems, the geometry of the rig and the correspondence between bar end faces across views, to establish in each case what the data determine and what remains ambiguous. Because the analysis separates the two, errors are not merely counted but traced back to their causes, in line with evaluation metrics that distinguish detection errors from association errors Luiten et al. 2021 [7].

The remaining errors. In the scenes on which the hybrid method fails completely, the incorrect association already yields the lowest cost. Consequently, a more accurate search algorithm alone cannot eliminate these errors. Further progress therefore requires evidence that distinguishes individual bars or a way to recognise the scenes at risk (section 6.4).

The hybrid method and the installed heuristic. The heuristic installed in the original system pairs the faces of each stacking level based on their order and accepts a pair only if it passes checks on torsion, length and tilt (section 2.4). The hybrid method significantly outperforms it, raising correspondence recall from 0.835 to 0.929 and reducing the scenes on which pairing fails completely from 14 to 5. The advantage grows with the other detector, since under two-pass suppression the recall of the installed heuristic drops to 0.317 whereas the hybrid method remains at 0.911. This collapse stems from the checks on angles and normals, as the same order-based pairing without checks remains stable. Based on the data and setup available in this work, even the pairing criterion of the original system, reproduced with its replay tool, performs worse than all the evaluated methods.

Order as a baseline. The order-based anchor without checks constitutes a strong baseline for this dataset. The hybrid method achieves a higher recall (0.929 against 0.872) and fails completely in fewer than half as many scenes (5 against 13). On the 96 independent bundles in the dataset (section 2.1), however, the difference in recall is not statistically significant because the hybrid method’s advantage stems entirely from the six scenes in which it accepts the crossing hypothesis (section 2.4) and on which the order-based anchor without checks pairs no bar correctly. In the remaining scenes, the two methods perform equivalently (section 5.3).

A reusable evaluation. The evaluation procedure developed in this thesis can be reused because the correspondence annotation is anchored to physical positions in the scene and can therefore score the output of any detector. Confidence intervals are computed by resampling whole bundles because errors occur one scene at a time and are not independent within a scene. The synthetic twin also makes it possible to vary pitch regularity, transport and view truncation independently on scenes where the identity of each bar is known by construction. Finally, every numerical value reported in this thesis is traceable to a stored result and those obtained on synthetic data can be reproduced from public data (section 6.3).

6.3 Limitations

The rig was still being commissioned during this work, so the number of available scenes is limited. The 108 scenes of the dataset represent 96 independent bundles distributed across four independent cross-sections. Given this sample size, the difference between the hybrid method and the order-based anchor without checks is not statistically significant, although the hybrid method correctly handles the scenes with crossing bars, where the order-based anchor without checks associates no bar correctly. Furthermore, the parameters of the shared cost model and two extensions of the hybrid method were tuned on this same dataset. While leave-one-out validation across profile families consistently places these parameters among the top performers, it also yields many equivalent alternatives that perform worse on the excluded scenes. Consequently, the available data is insufficient to definitively fix these parameters and a separate validation dataset would be required.

The controlled experiments rely on a synthetic twin that reproduces the input to the detector rather than the entire 3D point cloud. Consequently, the findings apply directly to the YZ projection and support the measurements on real data without fully replacing real-world validation. Additionally, the ambiguity sweep is restricted to a single profile family, the only one in which a bundle of five bars remains entirely within the scan window under a cross-view transport of three pitches. While this establishes whether the underlying mechanism depends on transport, it does not quantify its magnitude across other profile families.

The comparison with the pairing criterion of the original system is based on configuration choices made within this study, which may differ from those implemented in the active production environment. For this reason, the comparison does not directly reflect the performance of the full production line.

Finally, the raw point clouds cannot be made public. While results obtained with synthetic data are fully reproducible using public cross-section drawings, those based on real data require access to proprietary datasets. The absolute position of the carriages remains undetermined, as no physical reference in the scene provides an absolute constraint.

6.4 Future work

Detecting scenes at risk. The errors of the hybrid method occur only in scenes that contain detections without an annotated partner, where the counting term no longer favours the true association over a shifted one. Consequently, identifying a scene at risk is often more straightforward than correcting its pairing. When the two views contain a different number of detections, this risk condition can be inferred directly from the counts. Conversely, when both views yield the same count but frame different subsets of bars, detection counts alone are insufficient. In such cases, a potential indicator is the cost margin between the lowest-

cost association and its shifted alternatives (section 4.4), where a small margin would flag a high-risk scene to avoid incorrect automatic pairing.

Evidence that distinguishes each bar. When view truncation eliminates the advantage of the true association, the geometry of a regular bundle provides minimal guidance and the length check fails to detect the shift (section 4.1). Features specific to each bar, such as saw marks or the surface texture of the end face, could provide this evidence. A similar principle is used in forestry applications to match log end faces based on the structure of the wood Schraml et al. 2015 [12], Xie et al. 2025 [16]. It remains to be determined whether sawn metal end faces contain sufficient distinctive detail relative to the resolution of current industrial 3D scanners.

Identifiability during rig design. The identifiability analysis in section 3.2 was applied retroactively to an existing rig. If conducted during the initial design phase, this analysis would predict which parameters of a proposed geometry are determinable from sensor data and which remain ambiguous. Furthermore, the cost decomposition in section 4.2 would answer two design questions, namely how irregular the pitch of a bundle must be for the true and shifted associations to have distinct geometric costs and how wide the scan window must be for both views to frame the same bars. Because both analyses rely exclusively on geometry and the analytical cost model, they require no empirical data and can be conducted prior to physical construction.

Bibliography

- [1] Z. Cai et al. “MessyTable: Instance Association in Multiple Camera Views”. In: *Computer Vision – ECCV 2020*. Lecture Notes in Computer Science. 2020, pp. 1–16. DOI: 10.1007/978-3-030-58621-8_1.
- [2] M. A. Fischler and R. C. Bolles. “Random sample consensus: a paradigm for model fitting with applications to image analysis and automated cartography”. In: *Communications of the ACM* 24.6 (1981), pp. 381–395. DOI: 10.1145/358669.358692.
- [3] T. Hodaň et al. “BOP Challenge 2020 on 6D Object Localization”. In: *Computer Vision – ECCV 2020 Workshops*. 2020. arXiv: 2009.07378 [cs.CV].
- [4] S. Holm. “A Simple Sequentially Rejective Multiple Test Procedure”. In: *Scandinavian Journal of Statistics* 6 (1979), pp. 65–70.
- [5] H. W. Kuhn. “The Hungarian method for the assignment problem”. In: *Naval Research Logistics Quarterly* 2.1-2 (1955), pp. 83–97. DOI: 10.1002/nav.3800020109.
- [6] H. Law and J. Deng. “CornerNet: Detecting Objects as Paired Keypoints”. In: *Computer Vision – ECCV 2018*. Lecture Notes in Computer Science. Springer, 2018, pp. 765–781. DOI: 10.1007/978-3-030-01264-9_45.
- [7] J. Luiten et al. “HOTA: A Higher Order Metric for Evaluating Multi-Object Tracking”. In: *International Journal of Computer Vision* 129.2 (2021), pp. 548–578. DOI: 10.1007/s11263-020-01375-2.
- [8] J. Maye et al. “Online self-calibration for robotic systems”. In: *The International Journal of Robotics Research* 35.4 (2016), pp. 357–380. DOI: 10.1177/0278364915596232.
- [9] F. Nobre and C. R. Heckman. *FastCal: Robust Online Self-Calibration for Robotic Systems*. 2019. arXiv: 1902.10585 [cs.R0].
- [10] Y. Ohta and T. Kanade. “Stereo by Intra- and Inter-Scanline Search Using Dynamic Programming”. In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* PAMI-7.2 (1985), pp. 139–154. DOI: 10.1109/TPAMI.1985.4767639.
- [11] E. Ristani et al. “Performance Measures and a Data Set for Multi-Target, Multi-Camera Tracking”. In: *Computer Vision – ECCV 2016 Workshops*. 2016. arXiv: 1609.01775 [cs.CV].
- [12] R. Schraml et al. “Towards the applicability of biometric wood log traceability using digital log end images”. In: *Computers and Electronics in Agriculture* 119 (2015), pp. 112–122. DOI: 10.1016/j.compag.2015.10.003.

- [13] M. Shelukhan et al. *Ranking vs. Assignment: The Metric Mismatch in Multi-View Object Association*. 2026. arXiv: 2606.02022 [cs.CV].
- [14] Y.-H. Sun et al. *C-DOG: Multi-View Multi-instance Feature Association Using Connected δ -Overlap Graphs*. 2025. arXiv: 2507.14095 [cs.CV].
- [15] J. Unnikrishnan, S. Haghghatshoar, and M. Vetterli. “Unlabeled Sensing With Random Linear Measurements”. In: *IEEE Transactions on Information Theory* 64.5 (2018), pp. 3237–3253. DOI: 10.1109/TIT.2018.2809002.
- [16] Y. Xie et al. “Log End Face Feature Extraction and Matching Method Based on Swin Transformer V2”. In: *Forests* 16.1 (2025), p. 124. DOI: 10.3390/f16010124.
- [17] X. Zhou, D. Wang, and P. Krähenbühl. *Objects as Points*. 2019. arXiv: 1904.07850 [cs.CV].